



Truly Intelligent Circuit Design and Implementation

Zhiyao Xie

Dept. Electrical & Computer Engineering
Duke University

Outline of My Talk

- Part 1: My Ph.D. Works
- Part 2: My Future Plan

Electronic Devices are Everywhere



Designers Try to Deliver Generational Gains

iPhone 8, X

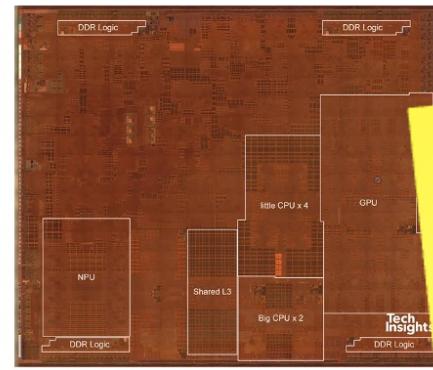


Apple A11

10nm

4.3 B trasistors

iPhone XS, XR

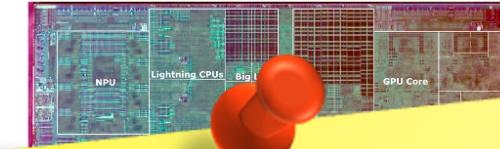


Apple A12

7nm

6.9 B trasistors

iPhone 11



Looks good!
Any challenges?

iPhone 12

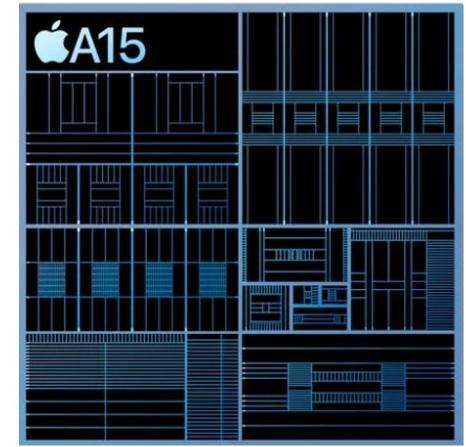


Apple A14

5nm

11.8 B trasistors

iPhone 13



Apple A15

5nm


15 B trasistors

Increased integration and architecture improvements

Chip Design Challenges

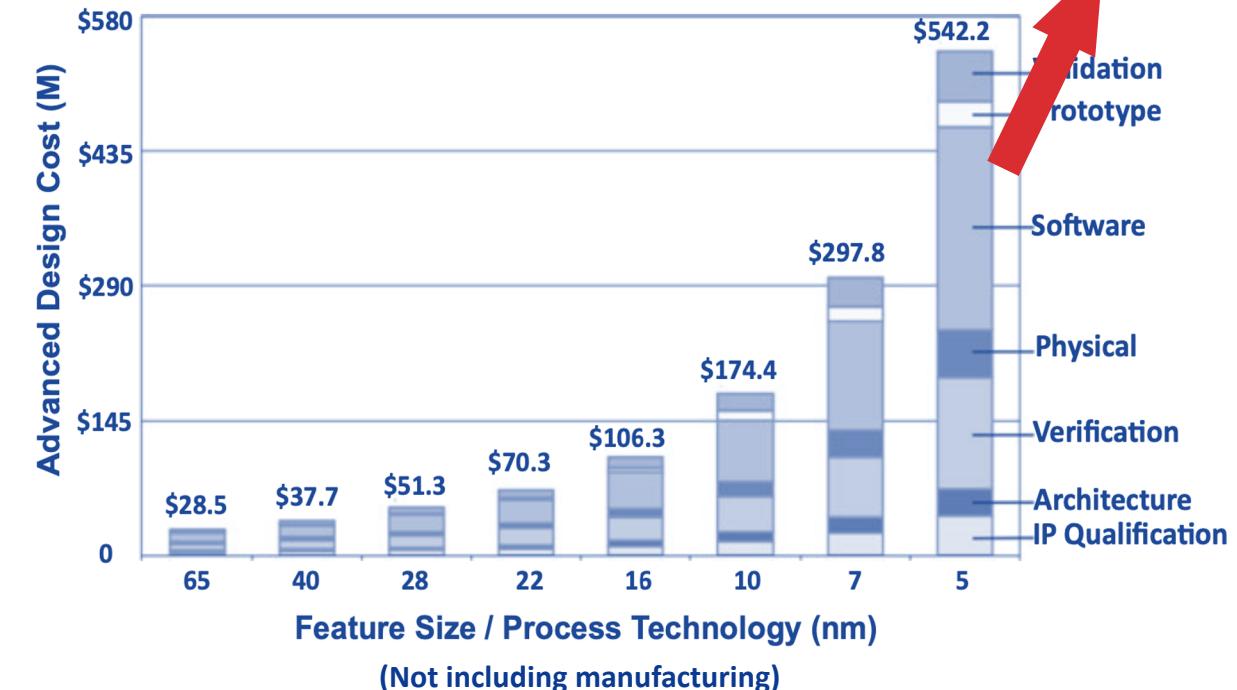
Diminishing performance gain and increasing design cost

Per-Core Performance Gain is Diminishing



Partially collected by M. Horowitz et al. Plotted by Karl Rupp, 2020

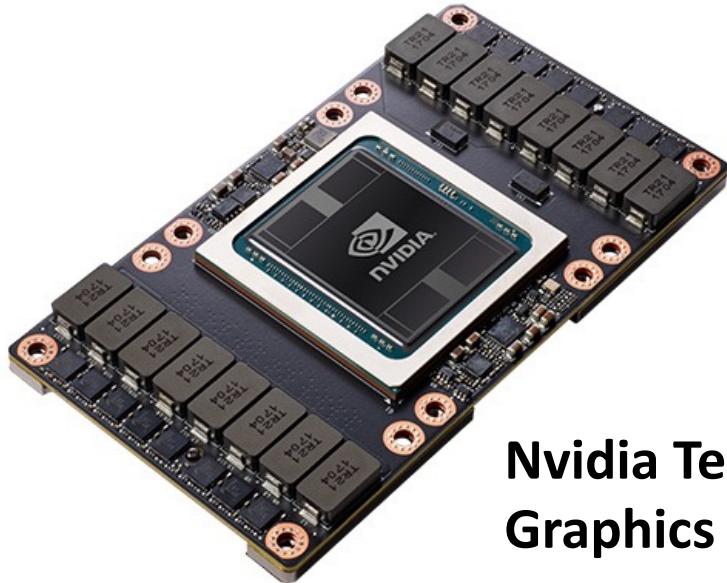
Design Cost is Skyrocketing



International Business Strategies, 2020

Chip Design Challenges

Not only costly, also long turn-around time



Nvidia Tesla V100
Graphics Card

It took **several thousand** engineers **several** years to create, at an approximate development cost of **\$3 billion**. – Jensen Huang, CEO of Nvidia

Nvidia GPU Technology Conference (GTC), 2017

This is Real Problem!

Challenges at advanced node

- Pressure from IPC and frequency
- Peak power keeps increasing
- Power delivery technique is
- Increasing design rules to me
- Increasing wire parasitics, ca
-
wire delay and noise

Intelligent design
methodologies
& solutions!



Inefficient chip design methodologies

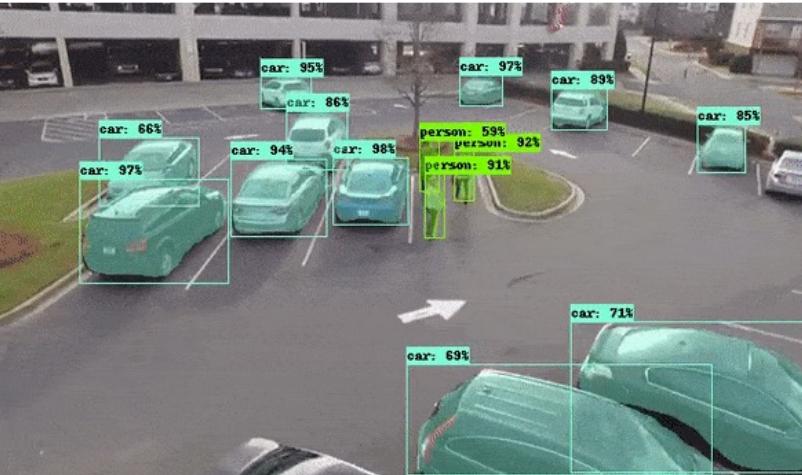


For one Arm CPU core
with ~3 million gates

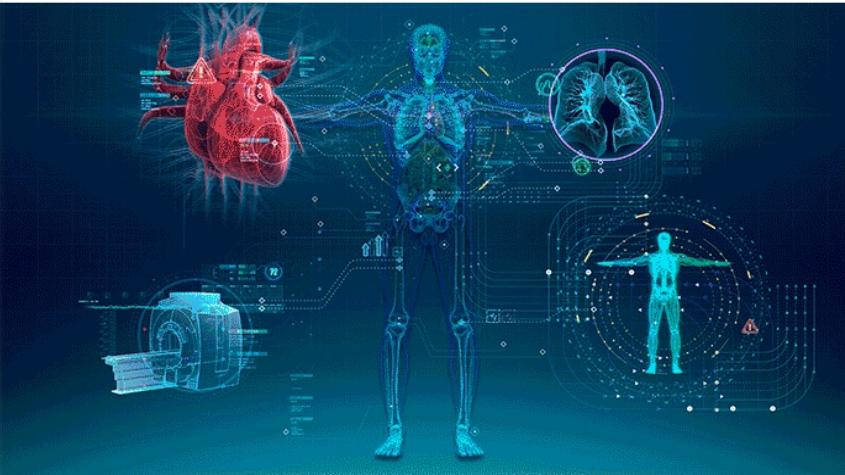
The power simulation takes ~2 weeks
Iteration in physical design take ~1 week
Solutions are repeatedly constructed from scratch
Solutions rely on designer intuition



https://github.com/ageitgey/face_recognition



<https://towardsdatascience.com/using-tensorflow-object-detection-to-do-pixel-wise-classification-702bf2605182>

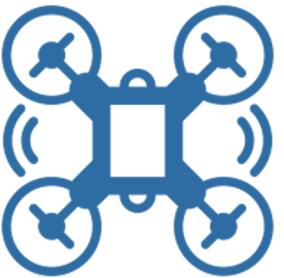


<http://matclinic.com/2017/05/18/the-team-behind-the-future-of-ai-in-healthcare/>

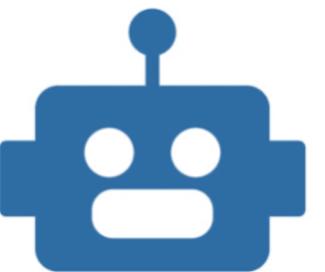
Self-driving Cars



Autopilot Drone



Robots



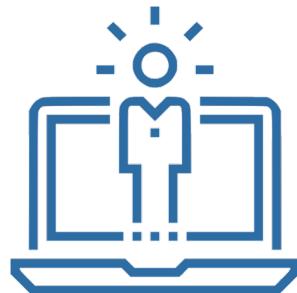
Smart Home



Health Monitor



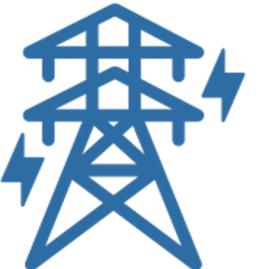
Personal Assistant



Manufacturing



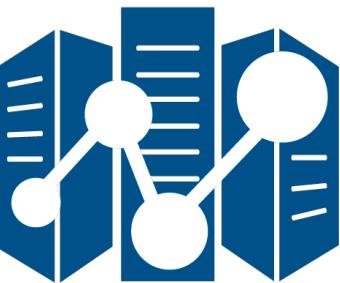
Smart Grid



Financial Service



HPC



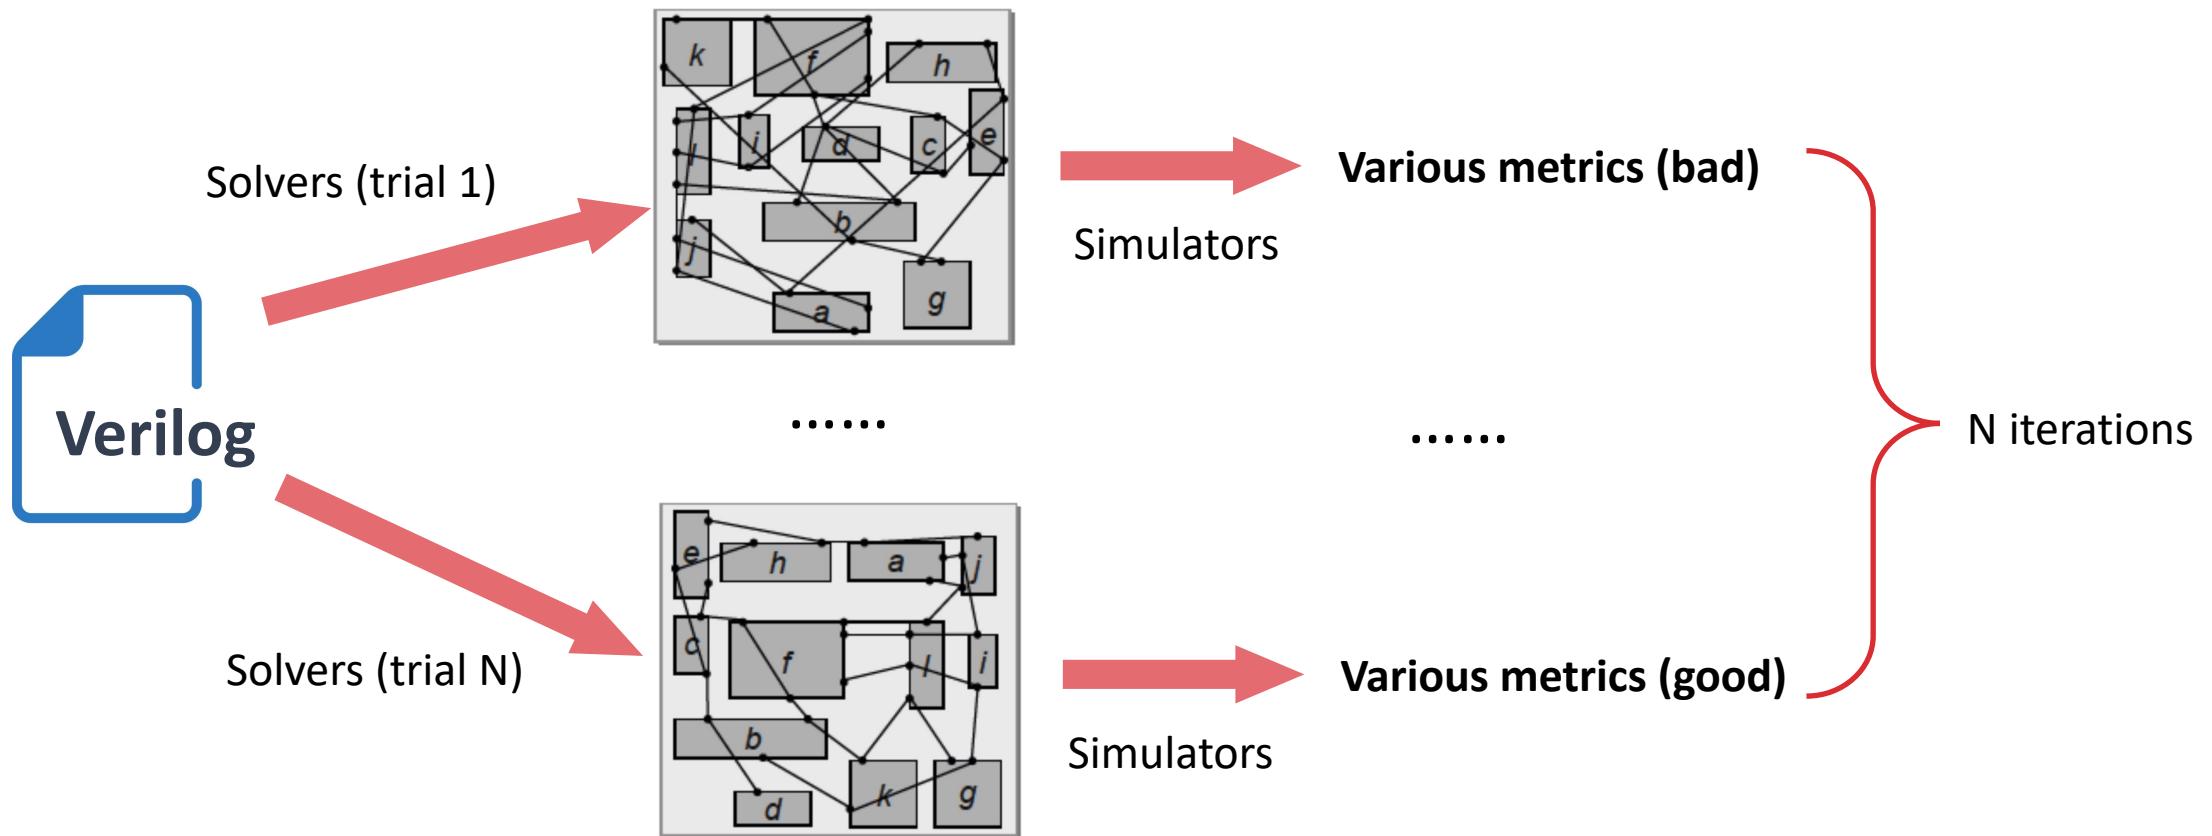
Security



Gaming

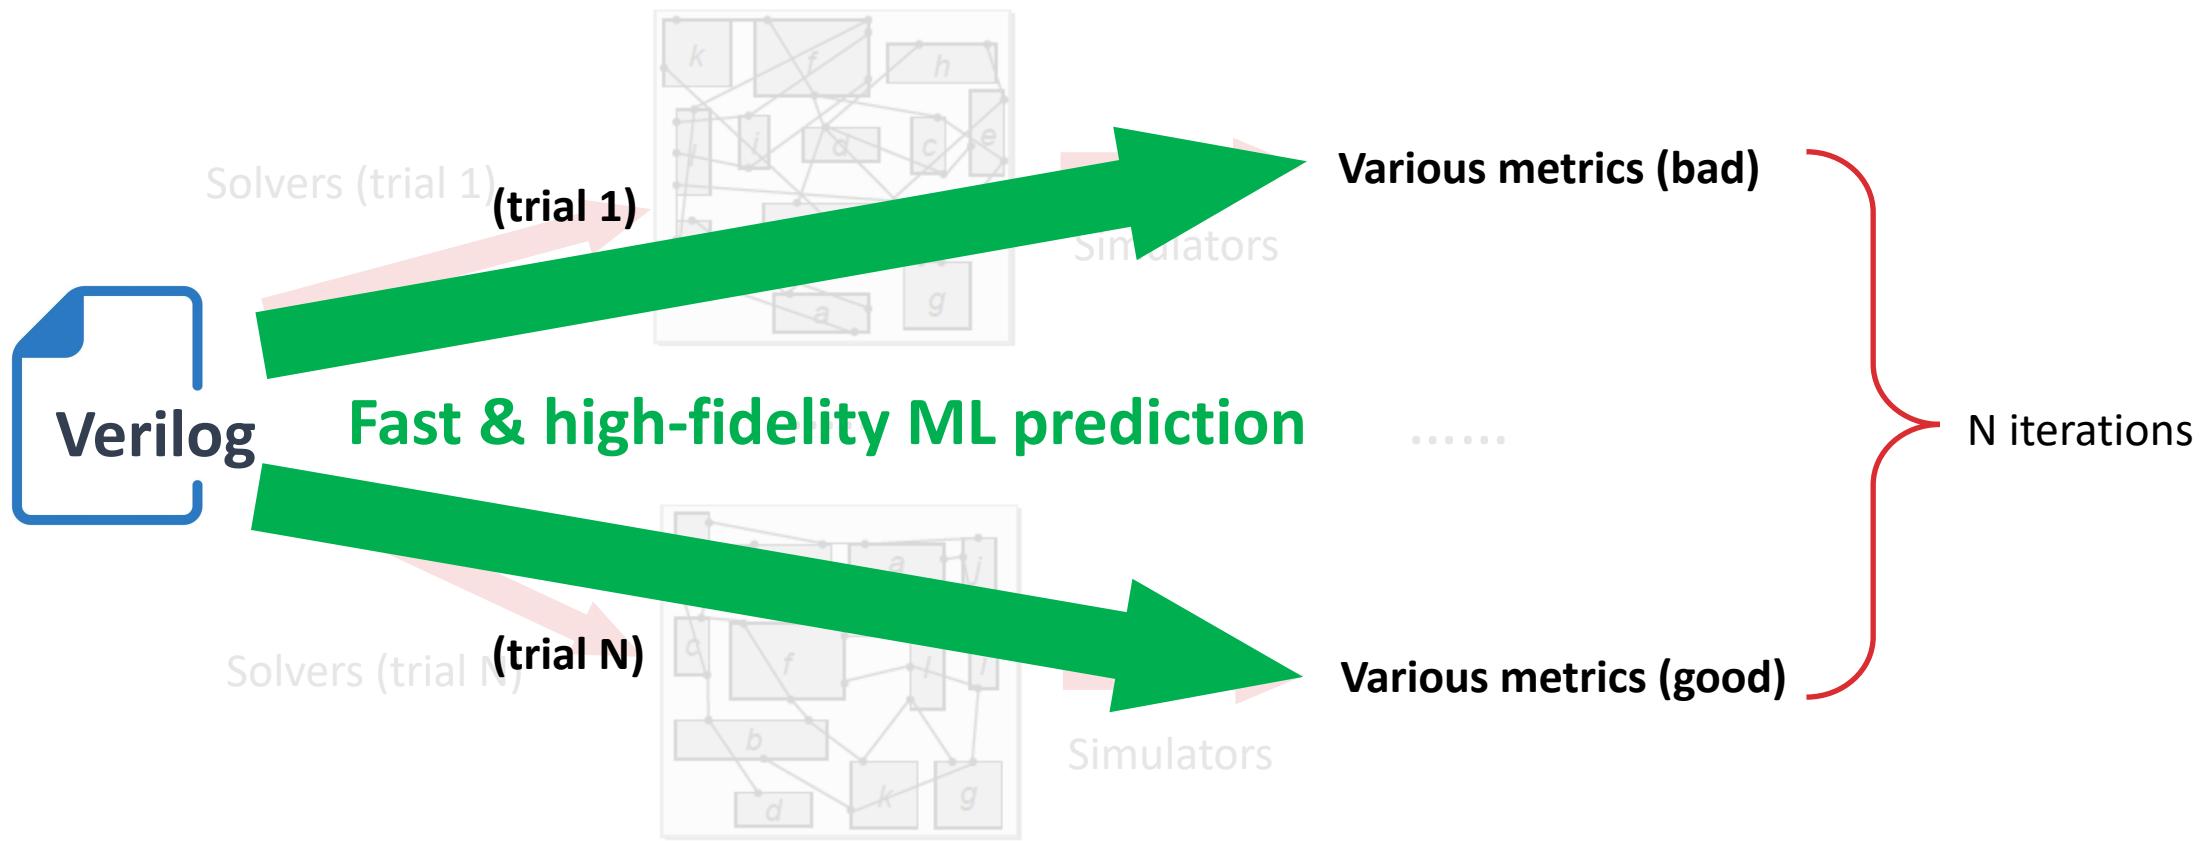


Why ML/Intelligence Helps Circuit Design?



- Producing solutions **repeatedly from scratch**

Why ML/Intelligence Helps Circuit Design?



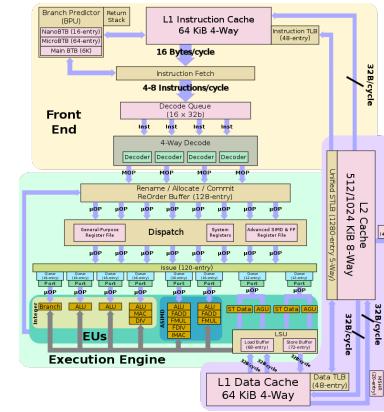
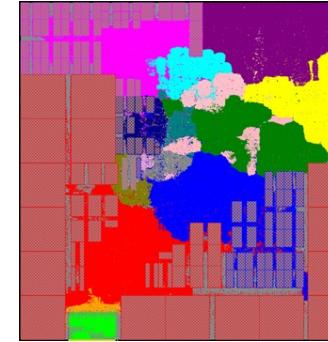
- Producing solutions **repeatedly from scratch**
- Why not learn from prior solutions?

Simple Plug-in and Use of ML Engines?



Images

- 100s * 100s pixels
- No extra information
- Any human can tell the label
- Data is everywhere
-



Circuits (Arm Neoverse N1 CPU core)

- Millions of connected components
- 100s GB of raw information
- Need simulations to get the label
- Data is hard to get

Innovative Customized Solutions are Desired!

Many Excellent Exploration in Academia and Industry


Increasing number of publications
on ML for chip design automation



UT Austin



UCSD



CUHK



TAMU



Duke



Cornell



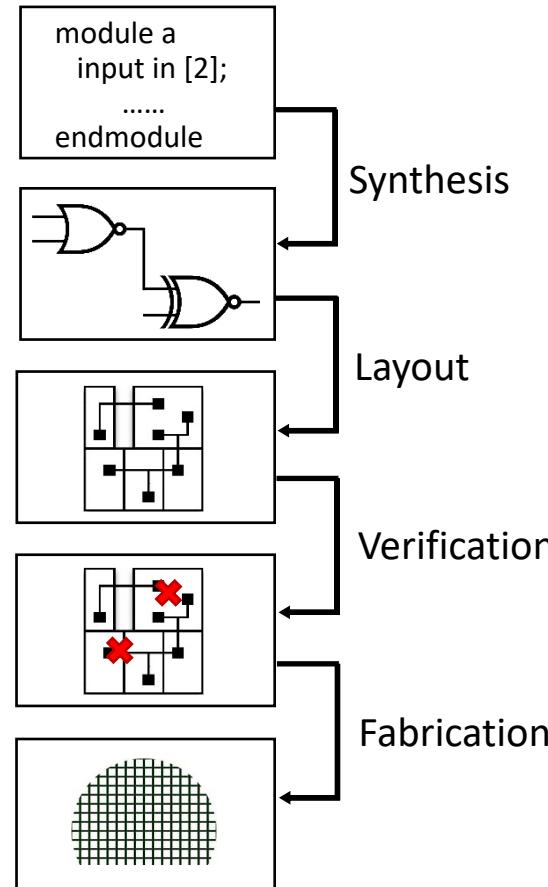
Google



Nvidia



GaTech



ML for EDA in
commercial tools

cadence

Cadence Innovus™

synopsys®

Synopsys ICCTM II

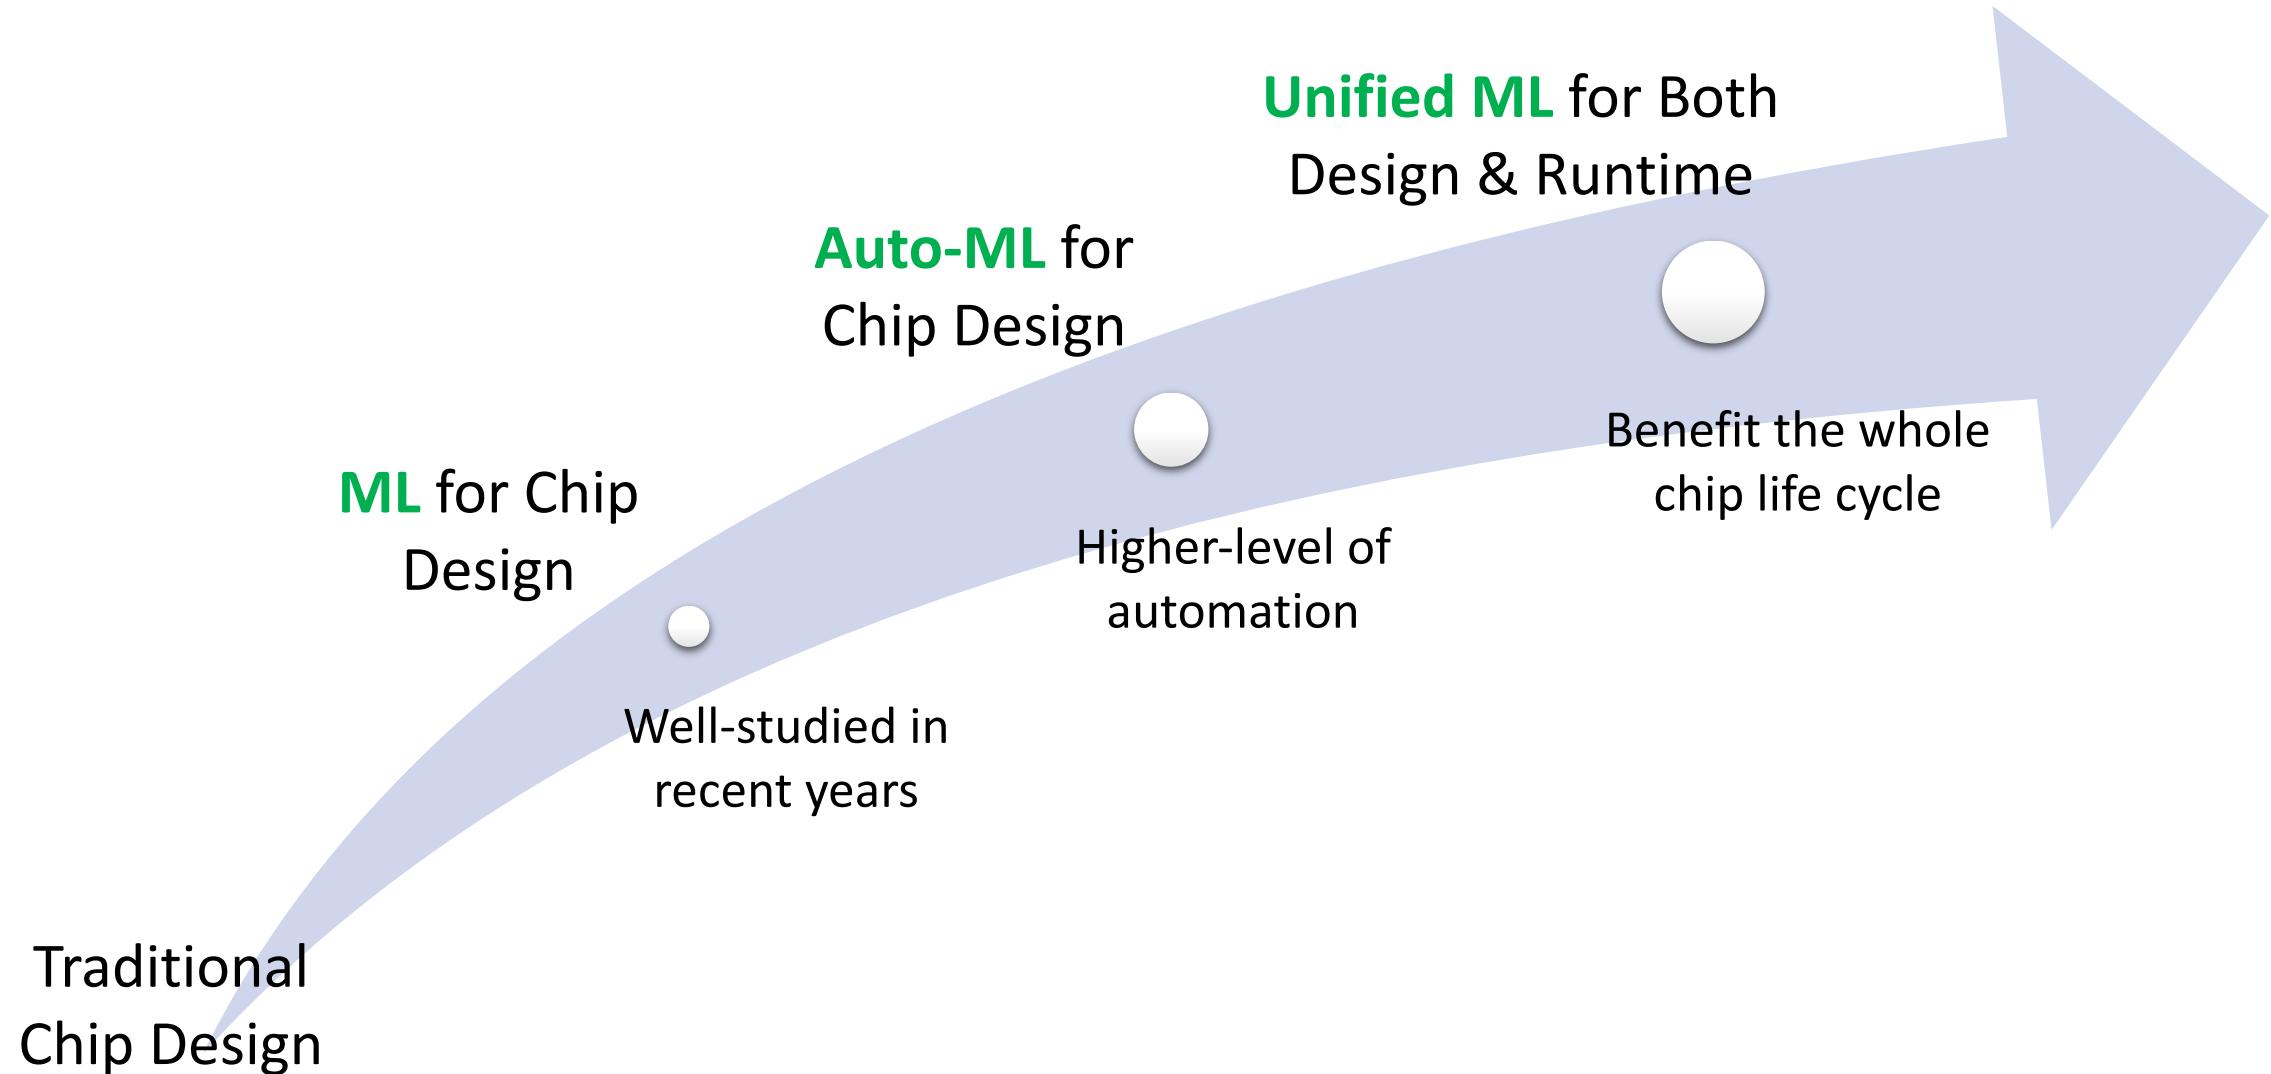
.....

ML for Chip Design



Electronics Research Initiative (ERI) – Design
Goal: 24 hours turnaround time & no human

What I Believe We Should Target



My Related Works

PPA

Power

Power & Power Delivery Challenges

[ICCAD'20], [ASPDAC'20],

[MICRO'21] (**Best Paper Award**)

Performance

Timing & Interconnect Challenges

[ICCAD'20], [ASPDAC'21],

[TCAD'21] (under review)

Area

Routability Challenges

[ICCAD'18], [DATE'18], [ICCAD'21]

Overall Flow Tuning

[ASPDAC'20]

Covered in this talk

My Related Works

PPA

Power

Power & Power Delivery Challenges

[ICCAD'20], [ASPDAC'20],

[MICRO'21] (**Best Paper Award**)

Performance

Timing & Interconnect Challenges

[ICCAD'20], [ASPDAC'21],

[TCAD'21] (under review)

Area

Routability Challenges

[ICCAD'18], [DATE'18], [ICCAD'21]

Overall Flow Tuning

[ASPDAC'20]

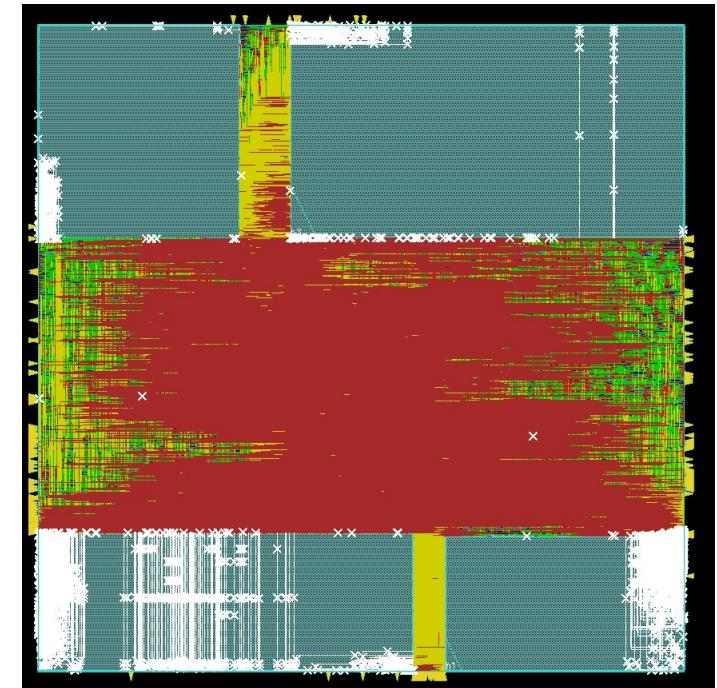
Covered in this talk

Case Study 1:

Routability Challenges

Routability Background

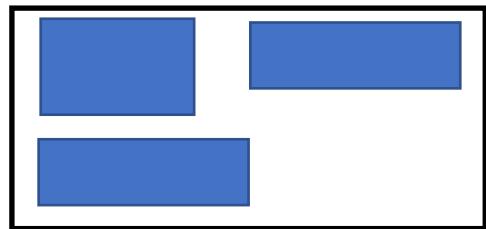
- Design Rule Checking (DRC)
 - Meeting manufacturing requirements
 - Less DRC violations (DRV) -> better routability
- DRV mitigation at early stages
 - Requires routability prediction/estimation
- Previous routability (DRV) estimations
 - ML model on small cropped regions
 - **Limited** receptive field and **missing** global information




DRC violations (white) on circuit layout

First Deep Learning Method for Routability Prediction

- Task 1: which one will result in less DRV count?



Layout 1




Layout 2




cat / dog

Customized CNN methods


- Task 2: where are DRC violations?



Layout 1



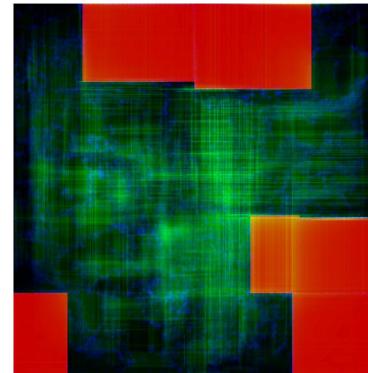
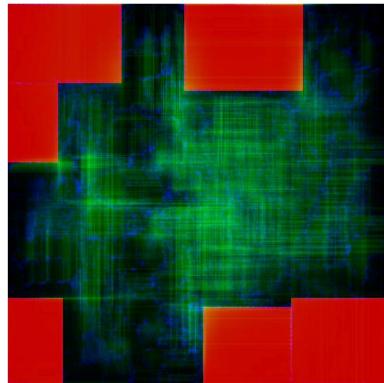
Layout 1



Customized FCN methods


First Deep Learning Method for Routability Prediction

- Task 1: which one will result in less DRV count?





- Requires global routing:
~~Hours * Number of Layouts~~
In seconds, with similar accuracy

- Task 2: where are DRC violations?




- Requires detailed routing
~~More hours * Iterations~~
In seconds, outperform previous works



Many Excellent Deep Learning Methods



J-Net [Liang, et al., ISPD'20]




PROS [Chen, et al., ICCAD'20]




Tremendous Engineering Efforts Required!

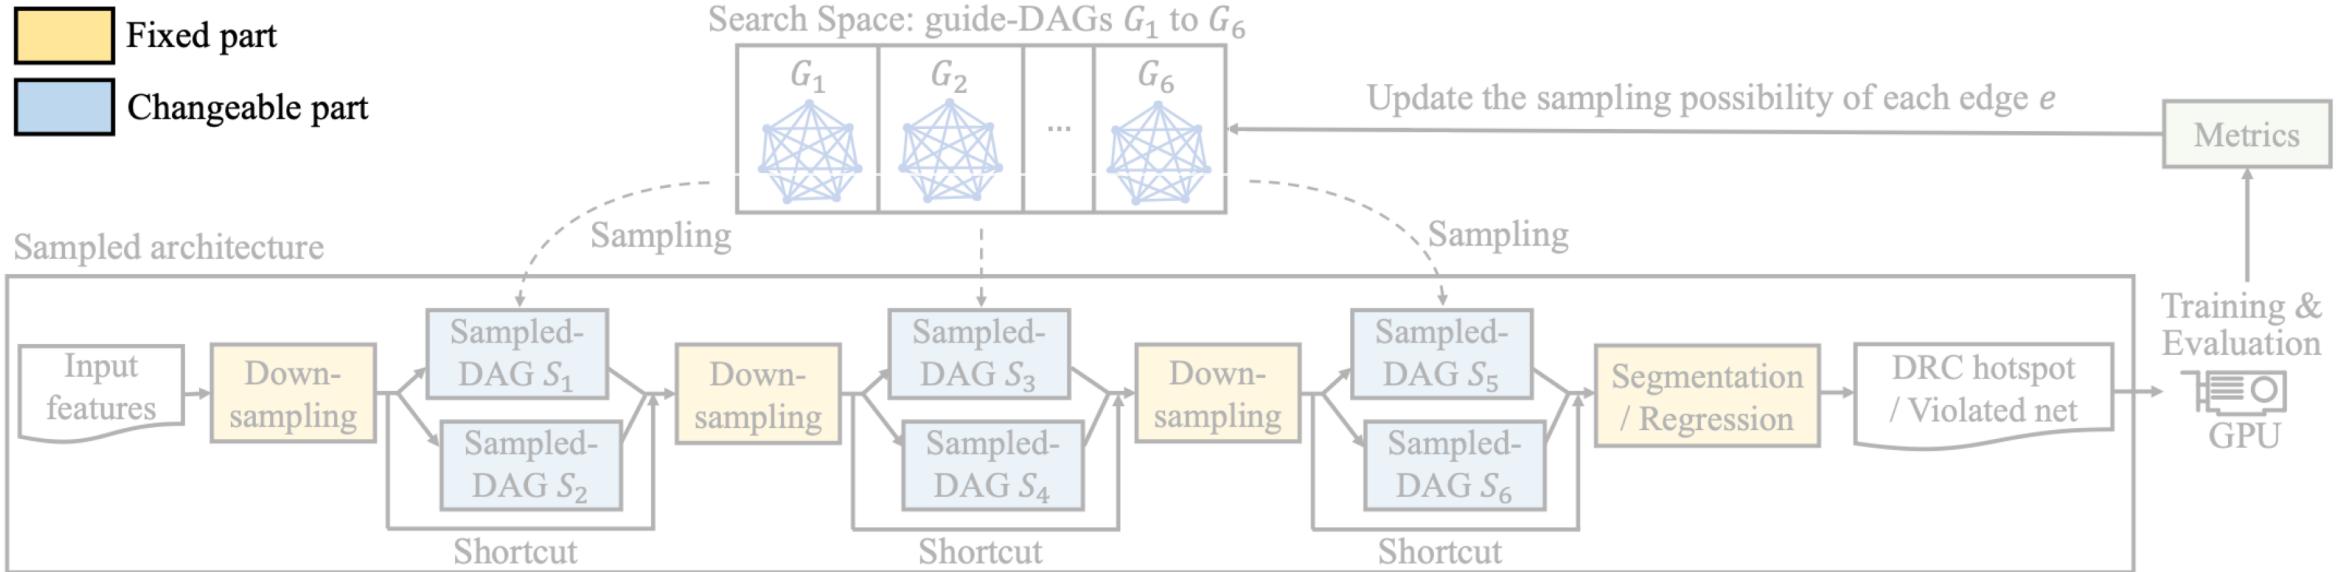
What I Believe We Should Target



Automatic Estimator Development – Search Space



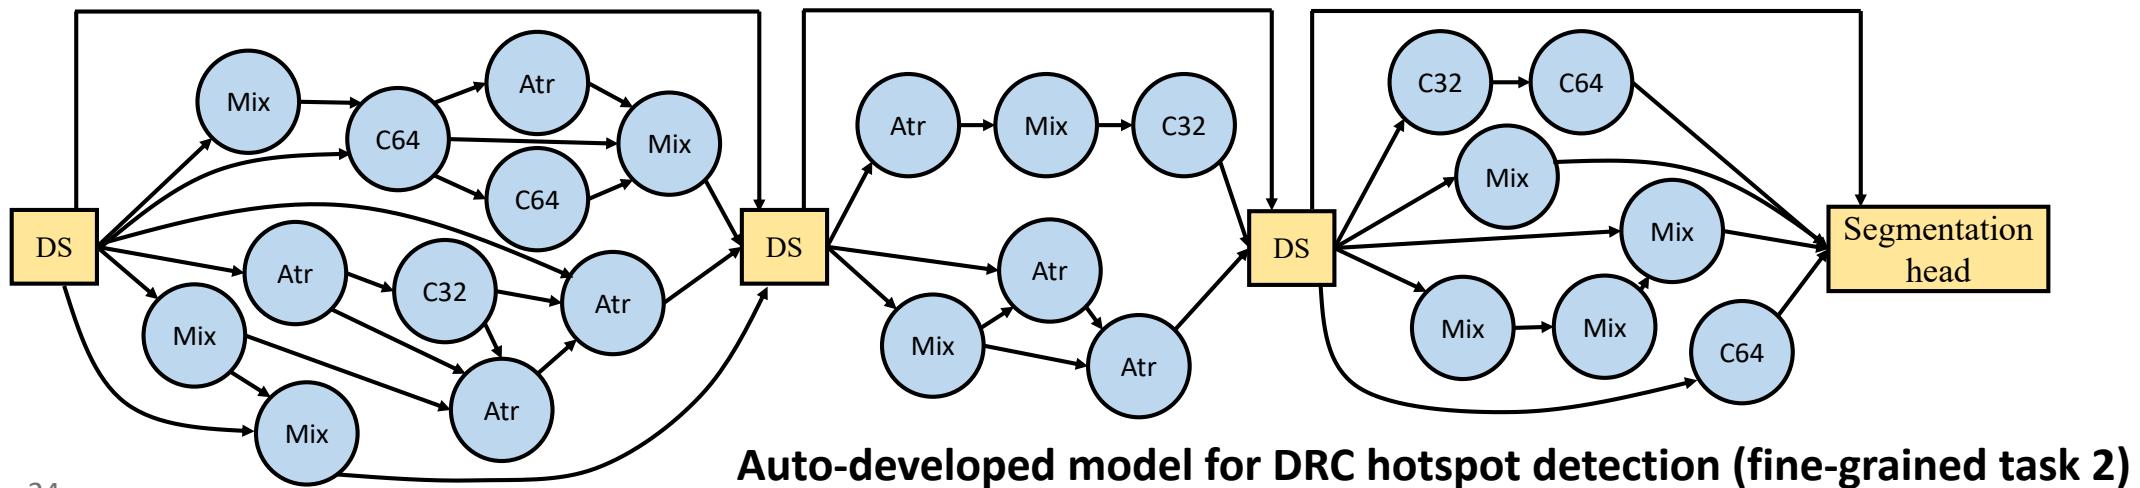
Automatic Estimator Development – Searching Algorithm



1. Sample from the completely-ordered graph (G_i) to get (S_i)
 2. Evaluate the sampled model by training and testing
 3. Update the sampling probability by evaluation result
- **Result: outperforms** previous works in both tasks; developed without human in **one day**


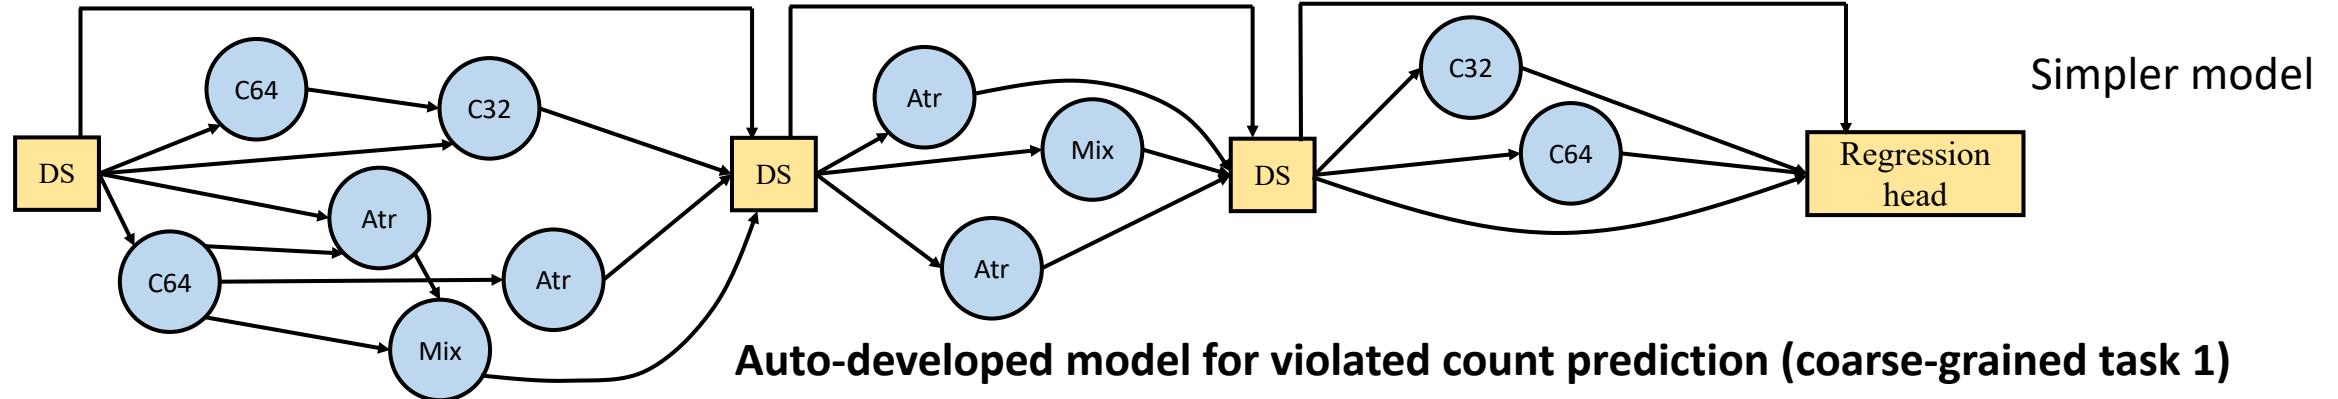
Auto-developed Model Structures

- Human-designed models:
 - Highly hierarchical and organized architecture
 - Limited operation types
- Auto-developed model:
 - Construct parallel branches and flexible interactions
 - Supports different operators



Auto-developed Model Structures

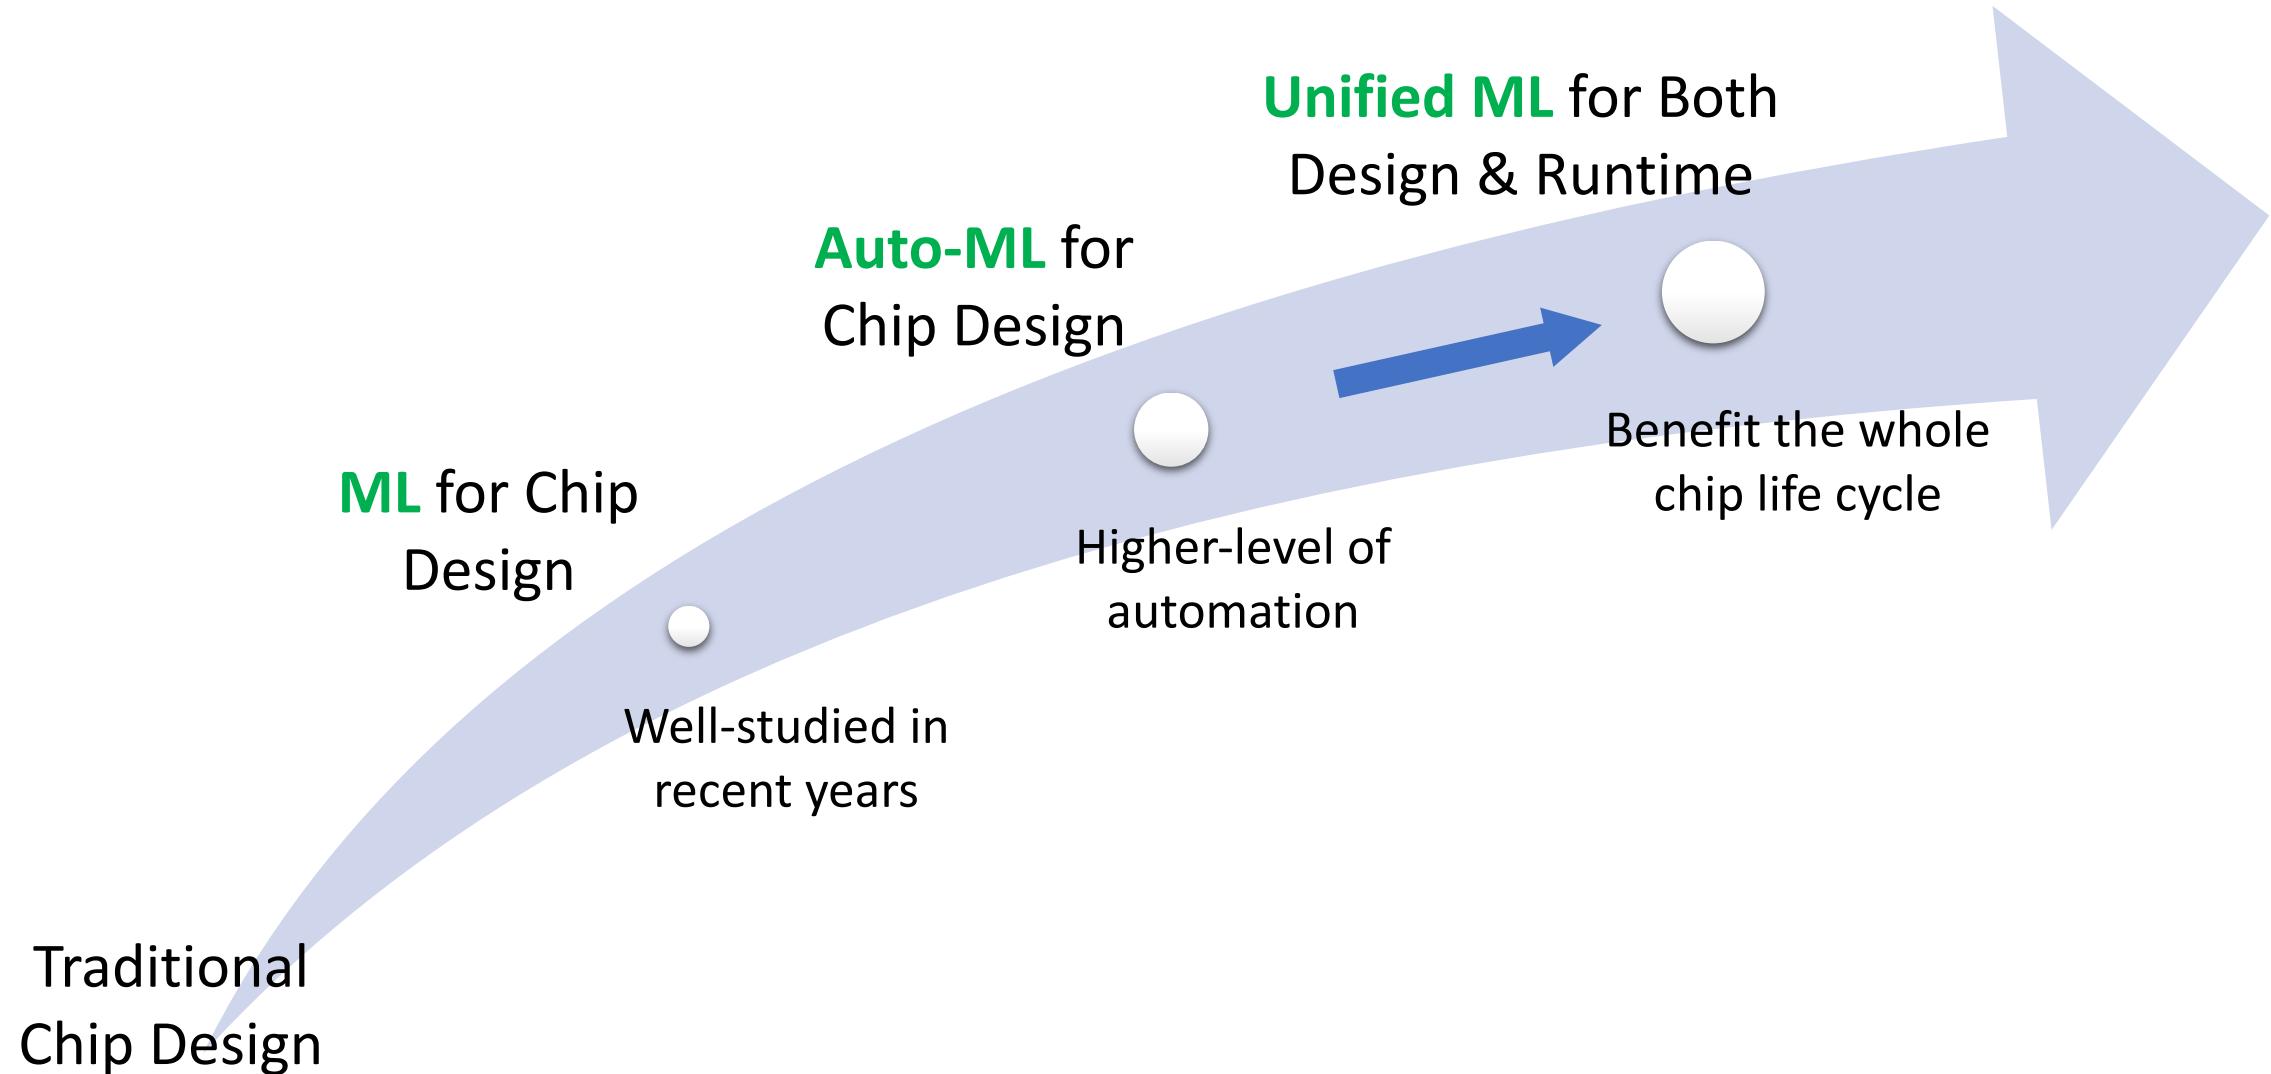
- Auto-developed model for DRC hotspot detection is significantly more complex




Case Study 2:

Power & Power Delivery Challenges


What I Believe We Should Target



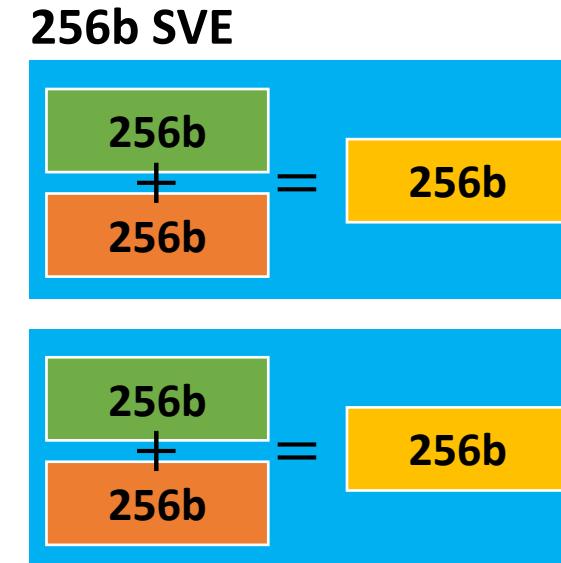
Challenge 1 – Design-time Power Introspection



Many-core CPU with
more transistors



Wider issue

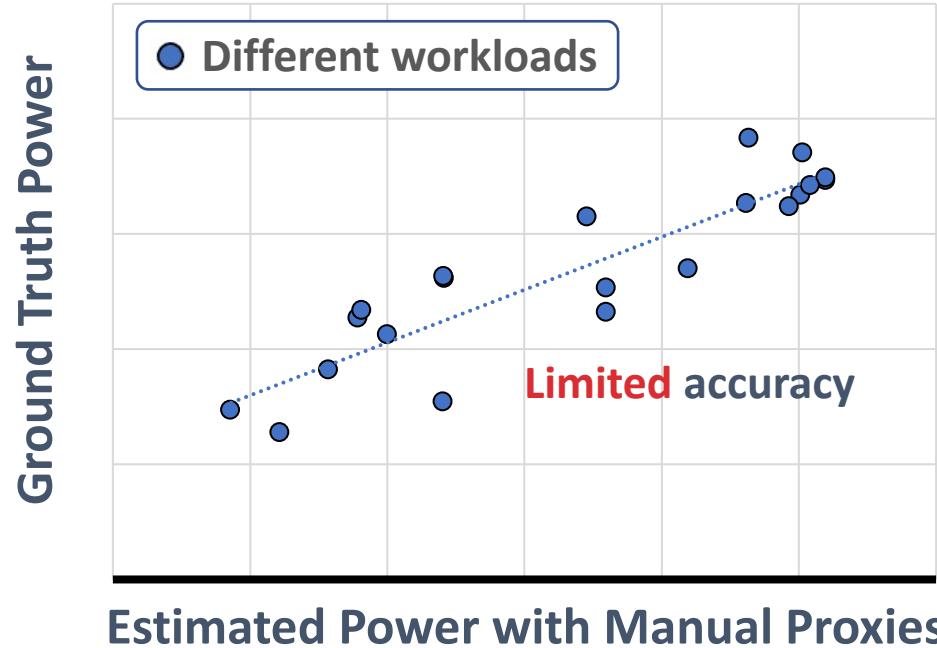


More vectored execution

- Delivering generational performance gains **adversely impacts** CPU power
- Power-delivery resources **not keeping pace** with CPU power demands
- **Increasing power-sensitivity** drives the need for design-time introspection

Challenge 2 – Run-time Power Introspection


Modelling power on one μarch block



- Peak-Power mitigation requires accurate power estimation to drive throttling


- Manually inferring proxies is very difficult in complex modern CPUs

- Abrupt changes in CPU current-demand (**di/dt event**) leading to deep voltage-droop



Challenge 3 - Workload Power Characterization

- Need power-characterization of **real-world** workloads
 - Simple micro-benchmarks not longer sufficient
- **Single SPEC simpoint can take weeks on the expensive emulator**
 - Power measurement is expensive
- **Only average power consumption available**
 - Impossible to scale to di/dt event analysis



Industry-Standard Emulator-Driven Power Flow

Challenges from Both Design-time and Runtime

A unified solution for both scenarios

Runtime Challenges Summary

- Peak power mitigation
 - **Difficult to manually** infer proxies
- Voltage droop (Ldi/dt) mitigation
 - Require very **low** response latency

Design-time Challenges Summary

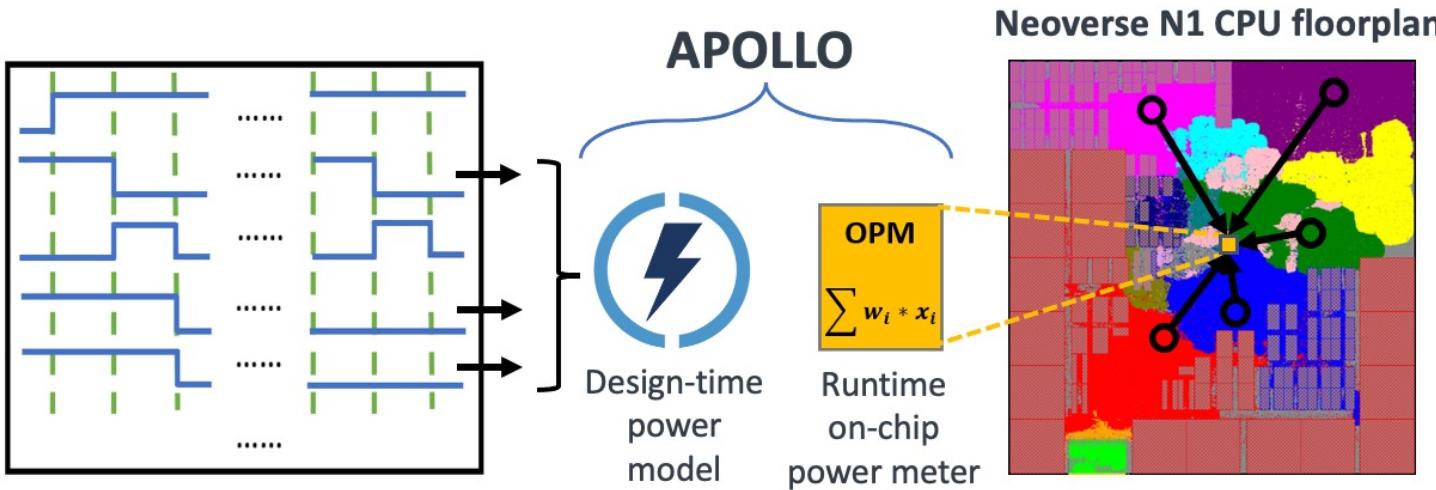
- Simulation on realistic workloads
 - **Expensive** and **slow**
 - **Limited** temporal-resolution

What is An “Ideal” Power Estimator?

1. Accurate yet fast
2. Achieve high temporal resolution
3. Low runtime on-chip overheads
4. Easily extensible to diverse designs




APOLLO: A Unified Power Modeling Framework



- **Fast, yet accurate design-time simulation**
- **Single-cycle** temporal resolution
- **Low-cost, yet accurate runtime monitoring**
- Design-agnostic **automated** development

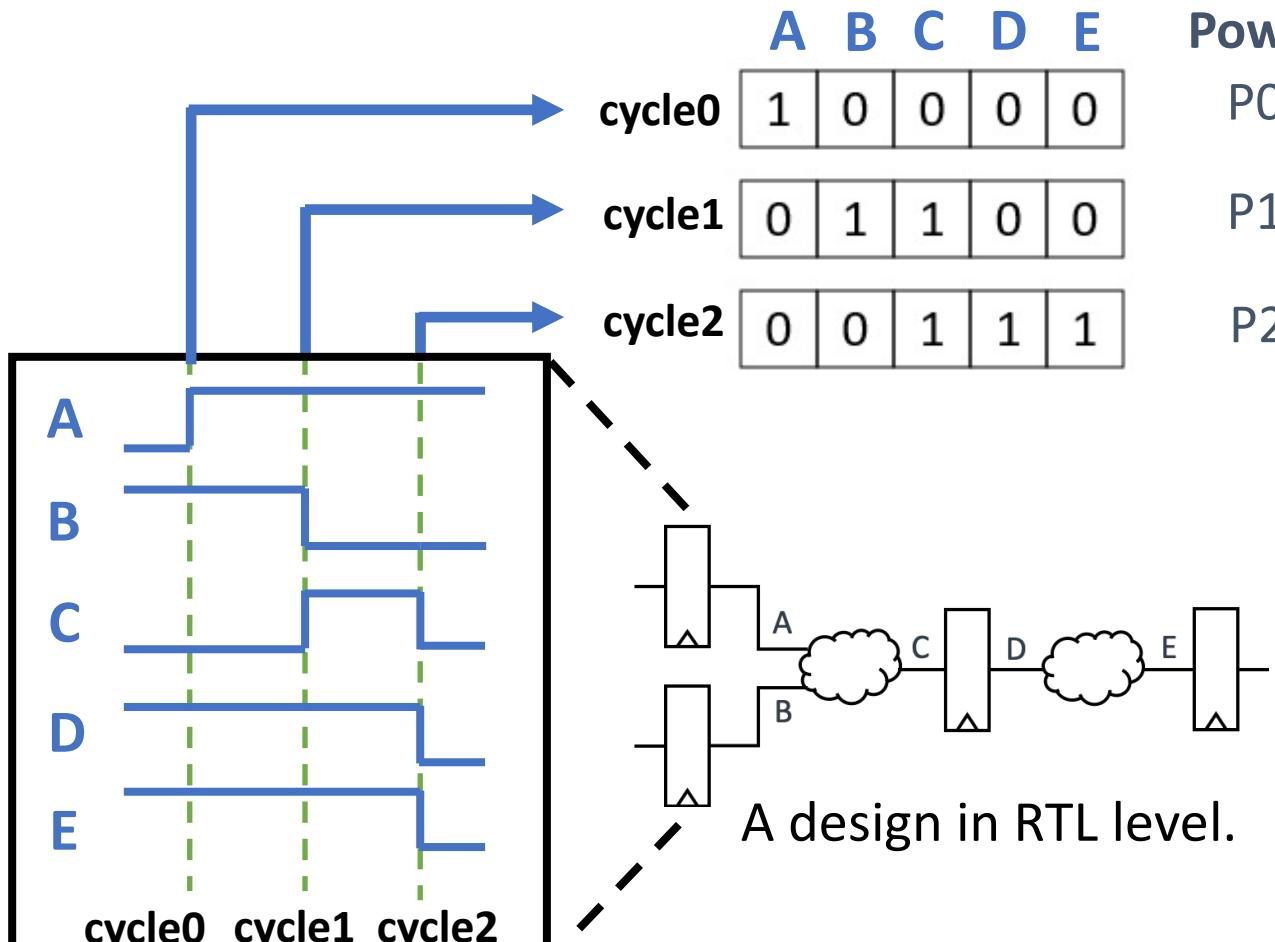
A Workload Execution Preview of APOLLO

40K cycles of APOLLO Power Estimation out of a trace of 17M cycles




- **~2 weeks execution time** reduced to **few minutes** on the emulator
- Unprecedented power-introspection due to **single-cycle** temporal resolution

APOLLO Feature Generation & Model Training



In .fsdb/.vcd file format

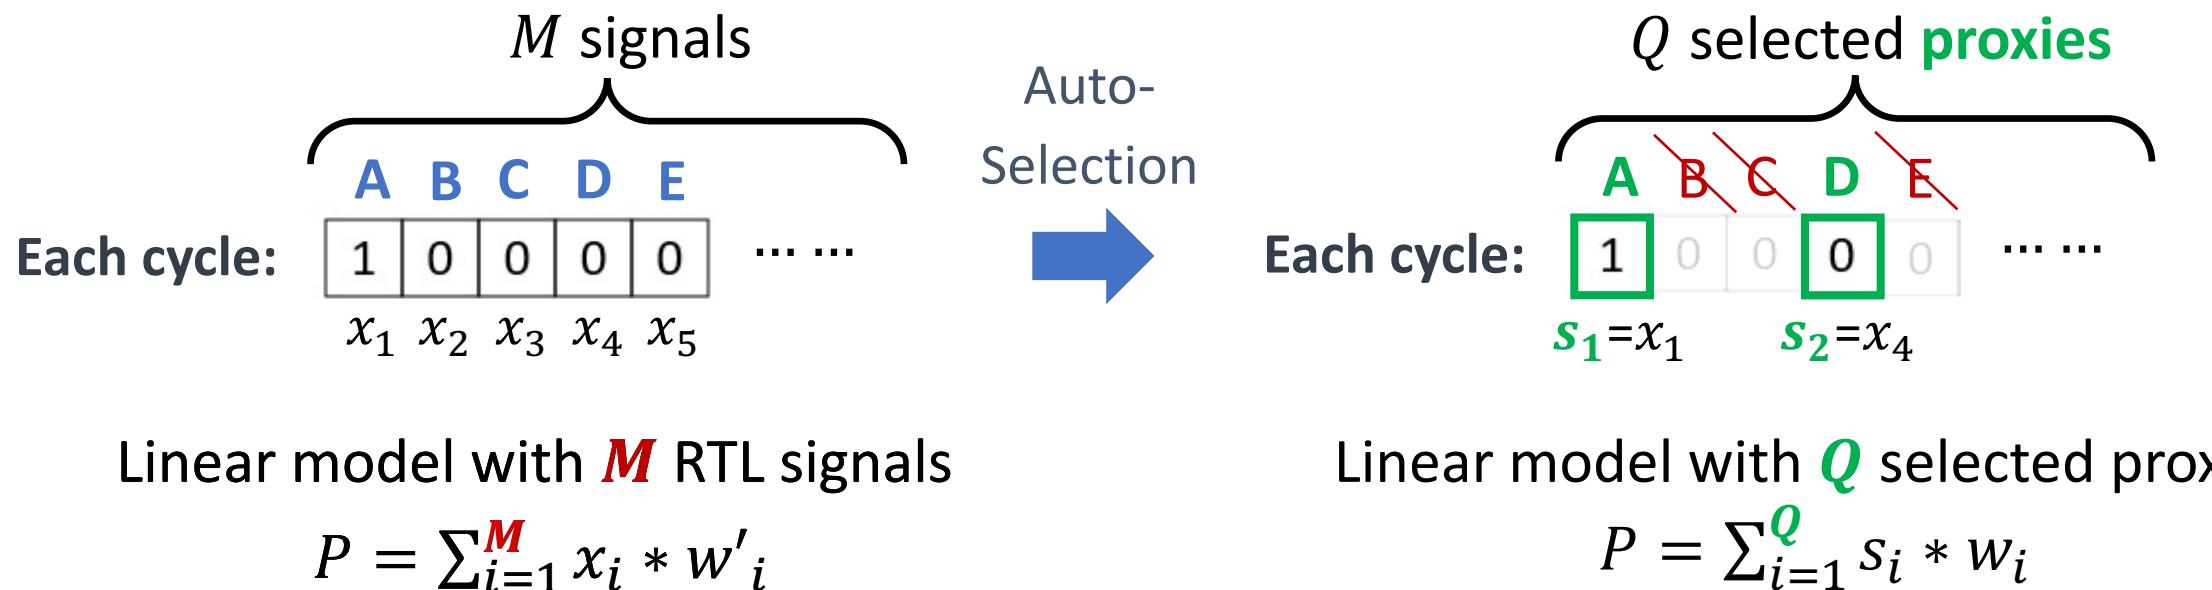
$M > 500,000$ in Neoverse N1
 $M > 1,000,000$ in Cortex-A77



Train the ML model: $F(X) = y$

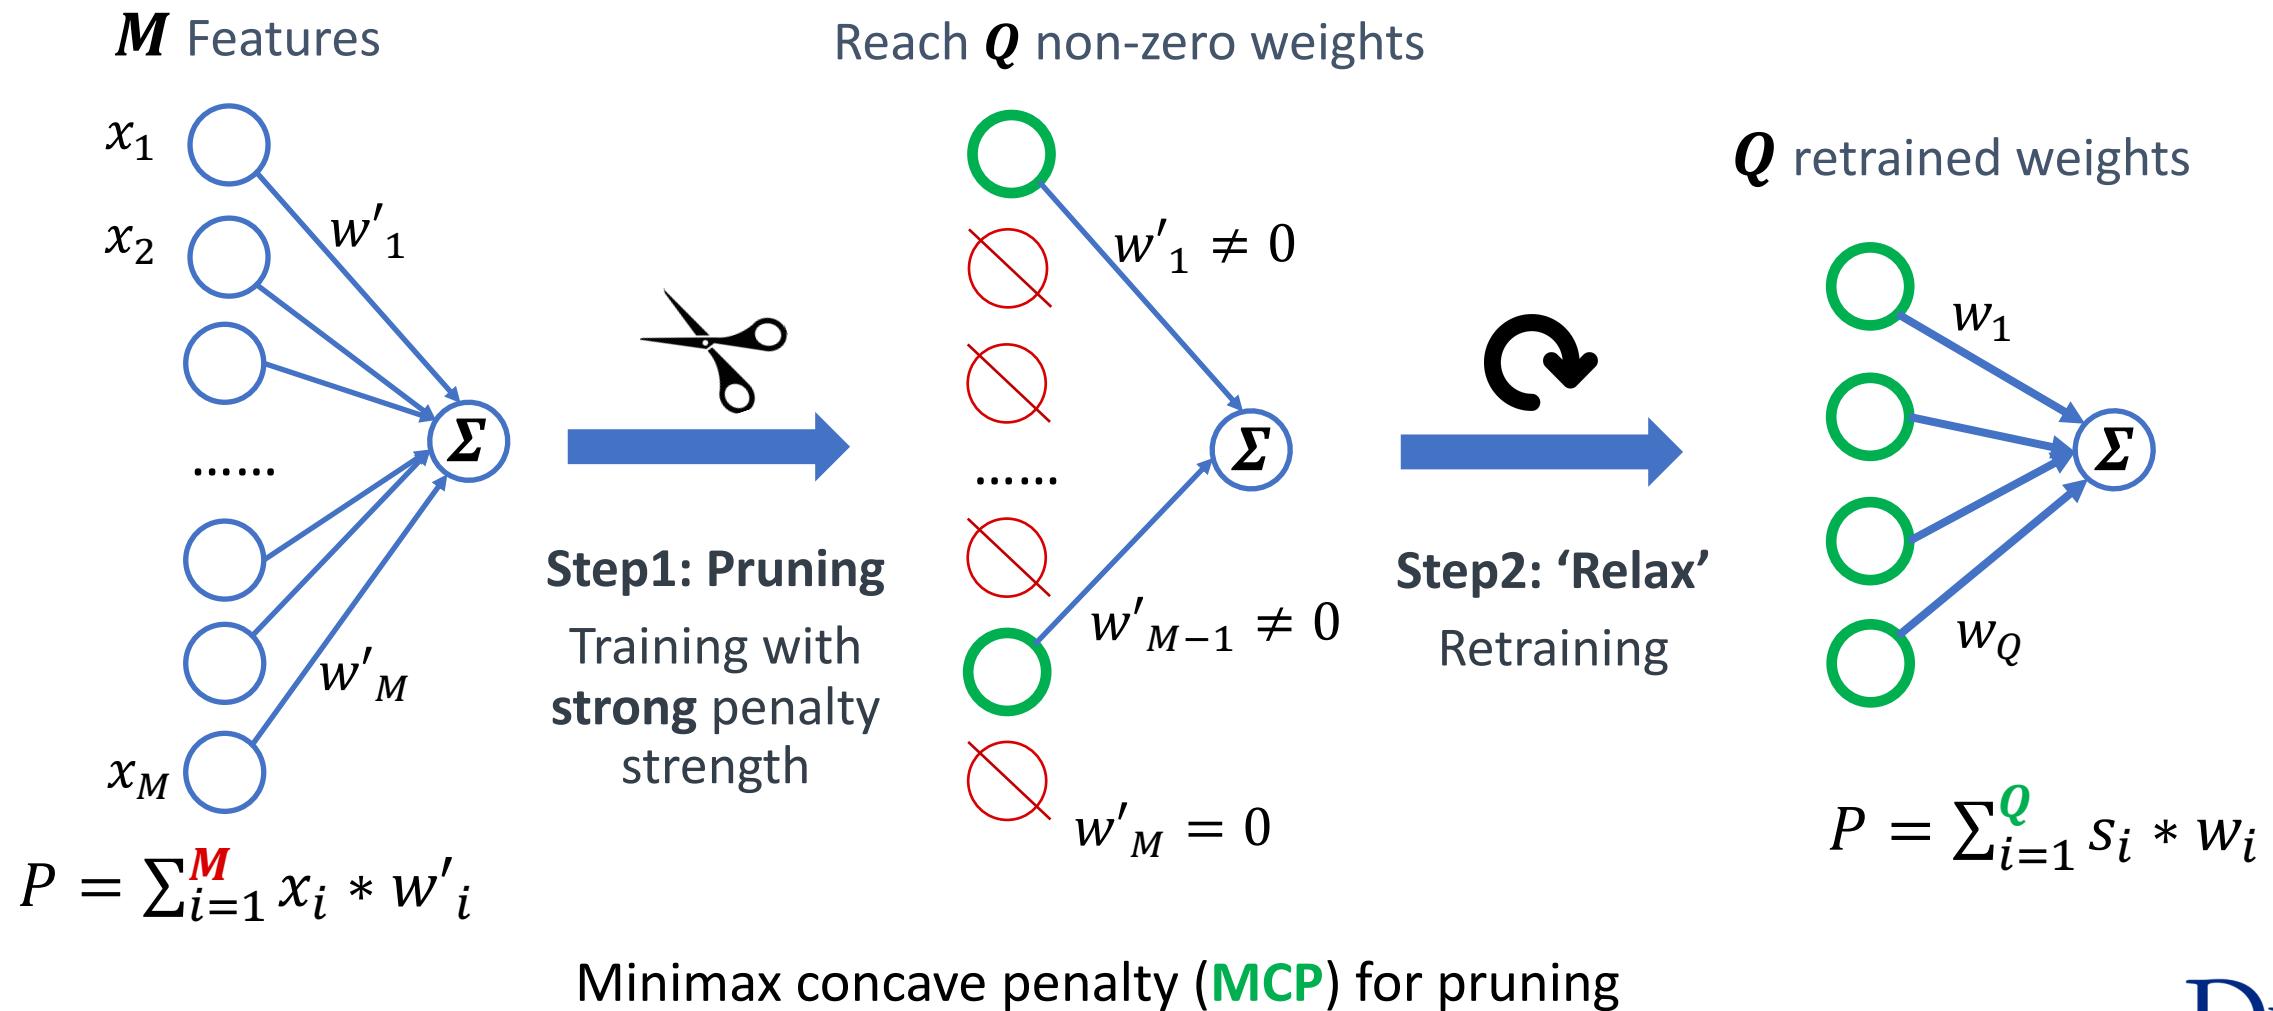
Simple Key Ideas

- **Linear** model can estimate power accurately
- **Small** portion of signals (proxies) can provide enough information



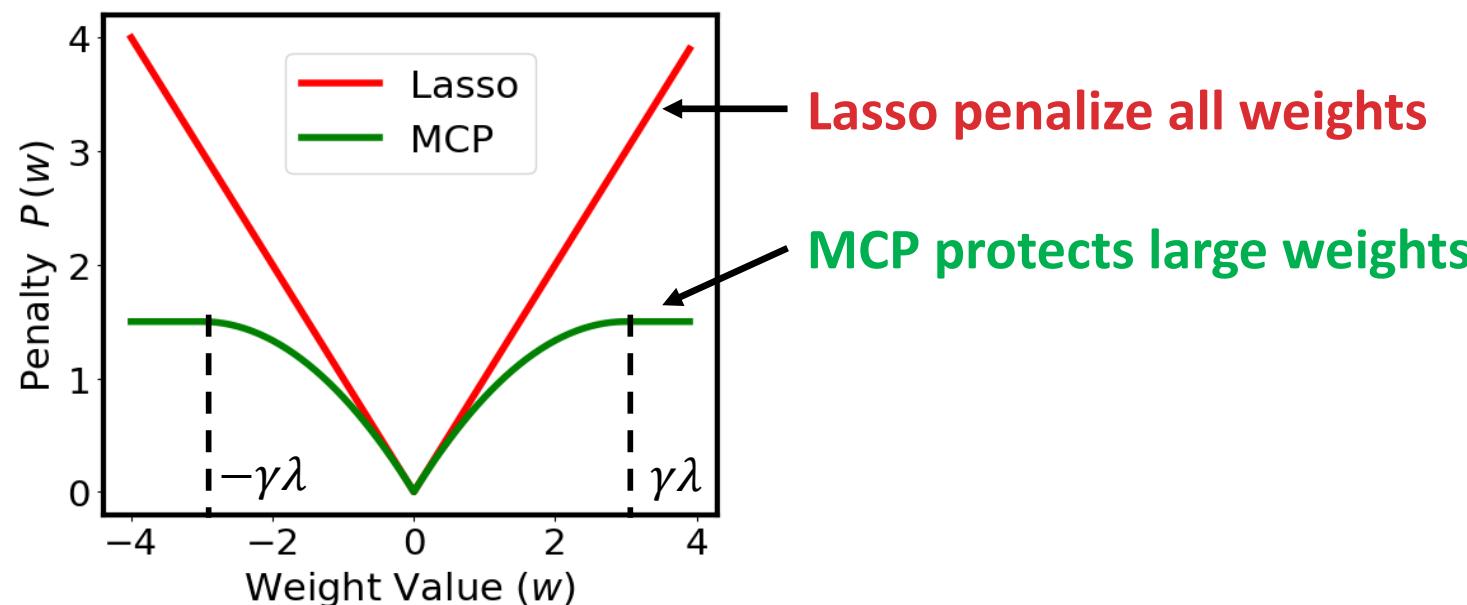
ML-Based Power Proxies Selection

Model construction in two steps



Why MCP for Pruning?

- To make $Q \ll M$, penalty is set to be very large.
- Lasso **degrades** model accuracy under large penalty
- MCP **protects** large weights thus **maintains** model accuracy




Model Training and Testing



Neoverse N1 (infra)
Deployed in AWS Graviton




Cortex A77 (mobile)
Deployed In Snapdragon 865



- Experiments on 3GHz 7nm Arm **commercial** microprocessors **Neoverse N1** and **Cortex A77**
- **Automatically** generate a “diverse” set of random micro-benchmarks for training
- Testing on **various** Arm power-indicative workloads

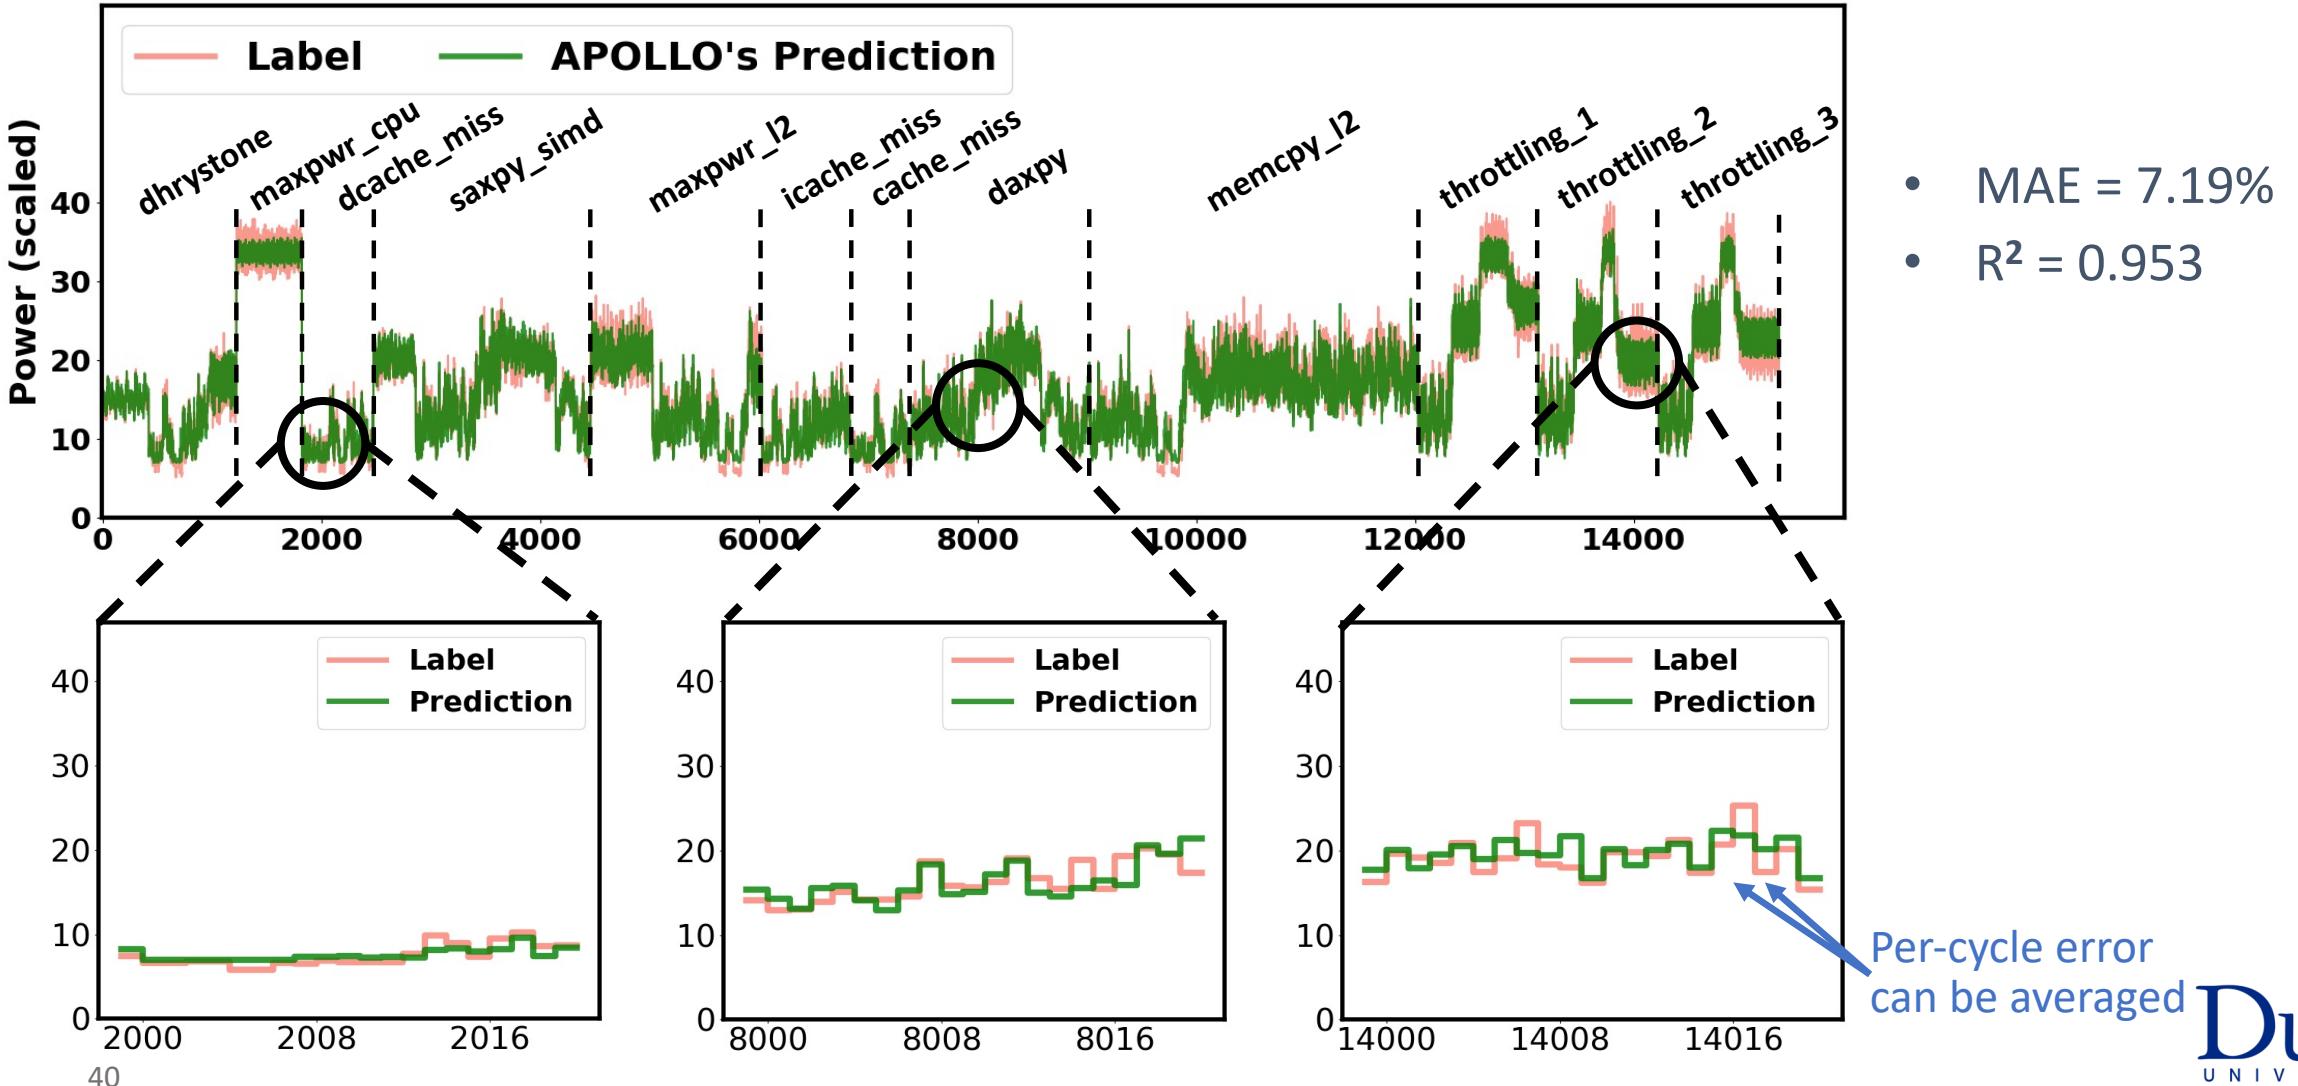
Prediction Accuracy as Design-Time Power Model

Per-cycle prediction from APOLLO with $Q=159$ proxies

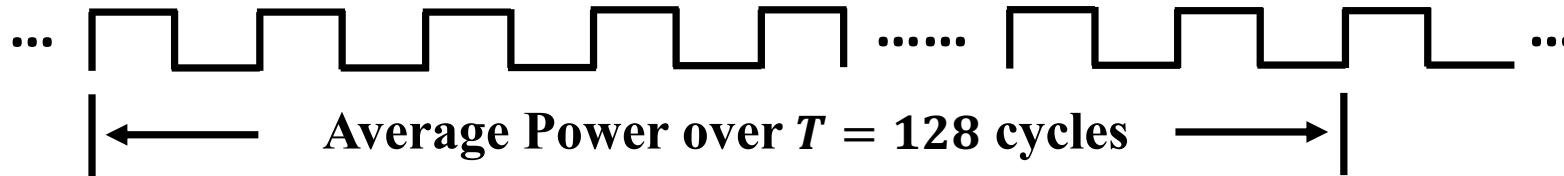


Prediction Accuracy as Design-Time Power Model


Per-cycle prediction from APOLLO with $Q=159$ proxies



Accuracy on Multi-Cycle Power Estimation



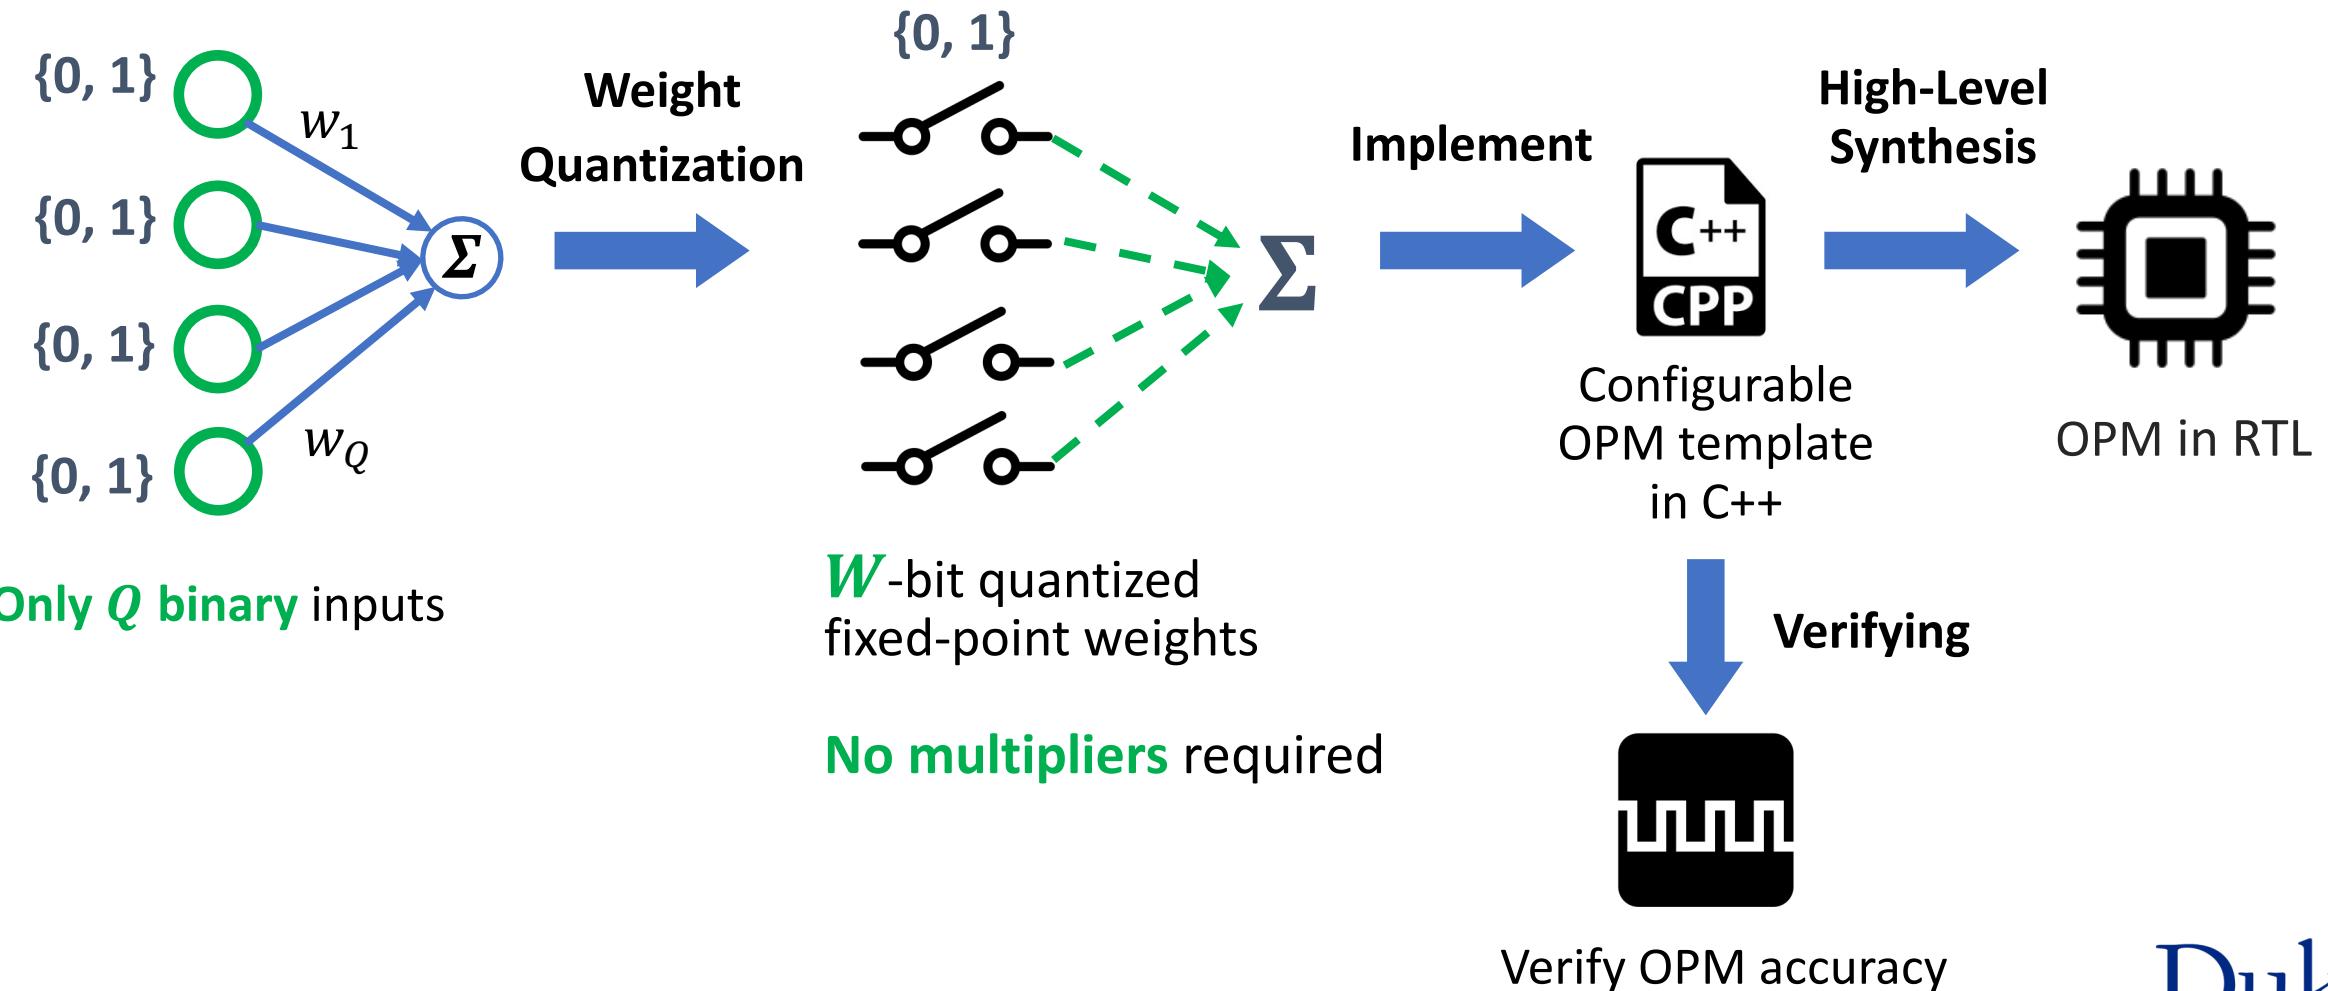
128-cycle prediction from APOLLO with $Q=70$ proxies



- MAE = 2.82%
- $R^2 = 0.993$
- Higher accuracy

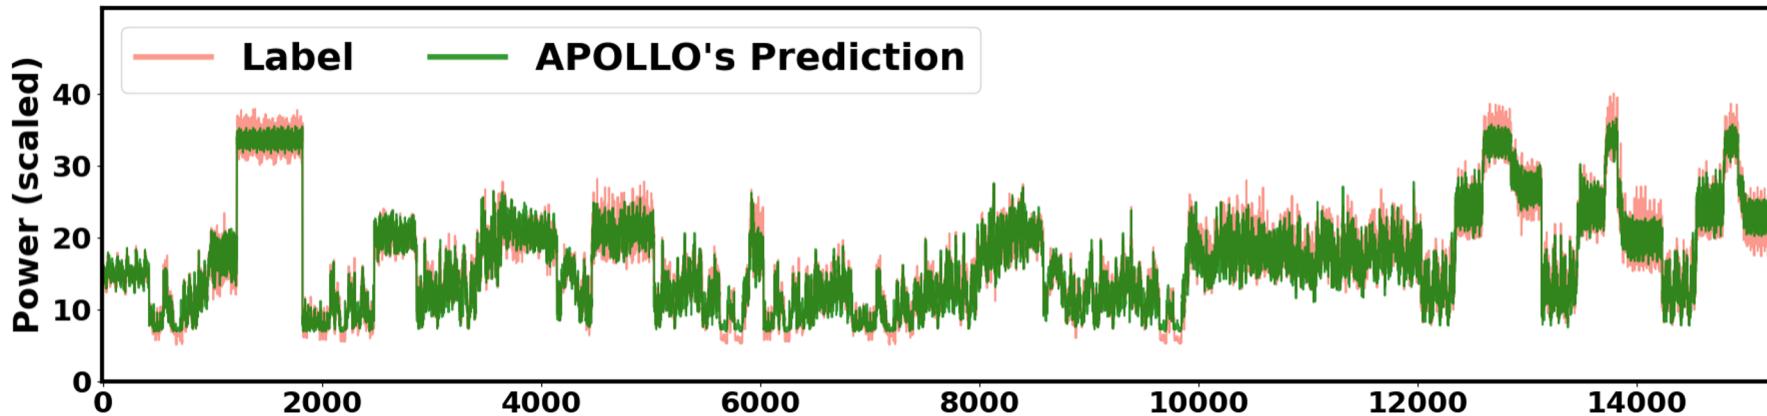
Automated Low-Cost Runtime OPM Implementation

APOLO is designed to be hardware-friendly



Prediction Accuracy from Design-time Model & OPM

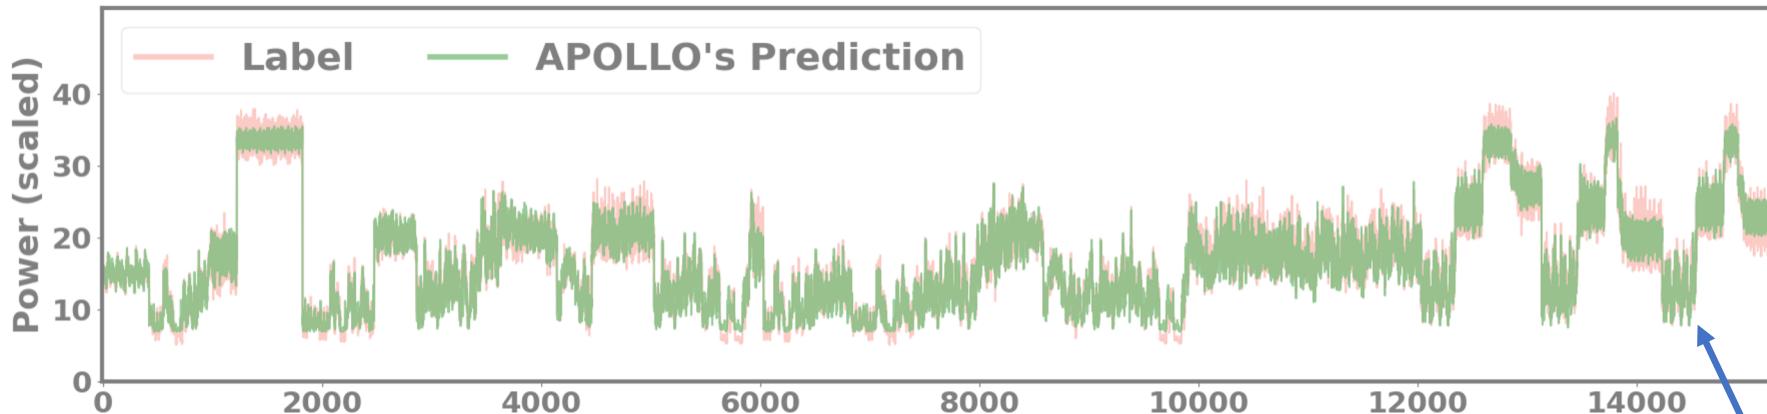
Per-cycle prediction from APOLLO with $Q=159$ proxies



- MAE = 7.19%
- $R^2 = 0.953$

Prediction Accuracy from Design-time Model & OPM

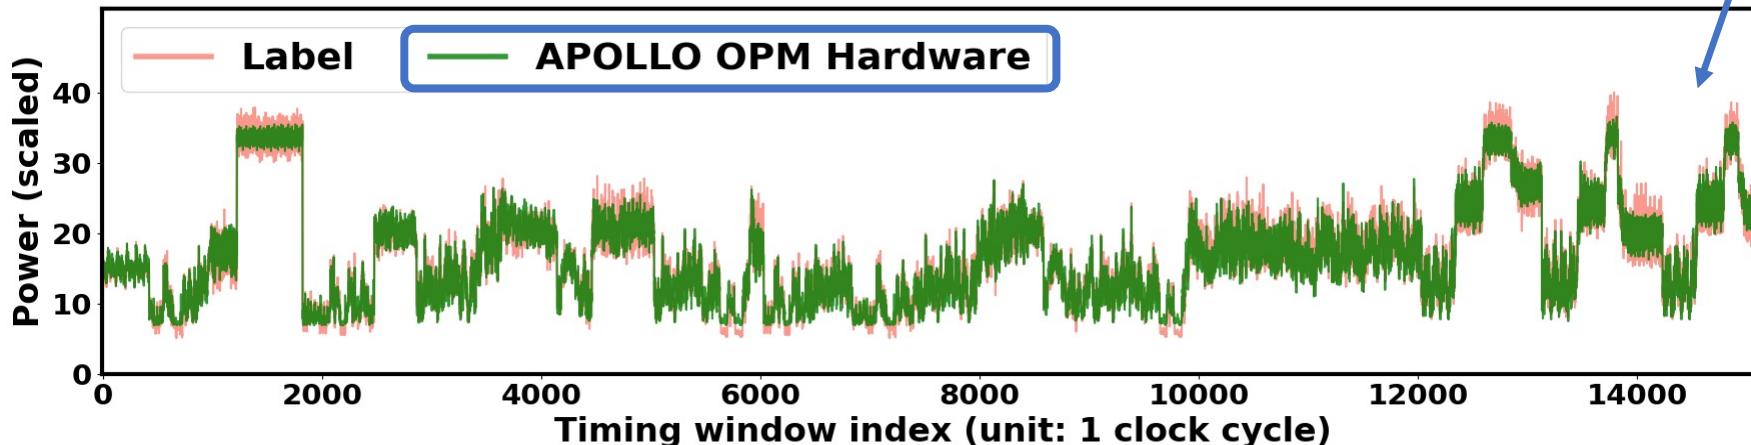
Per-cycle prediction from APOLLO with $Q=159$ proxies



- MAE = 7.19%
- $R^2 = 0.953$

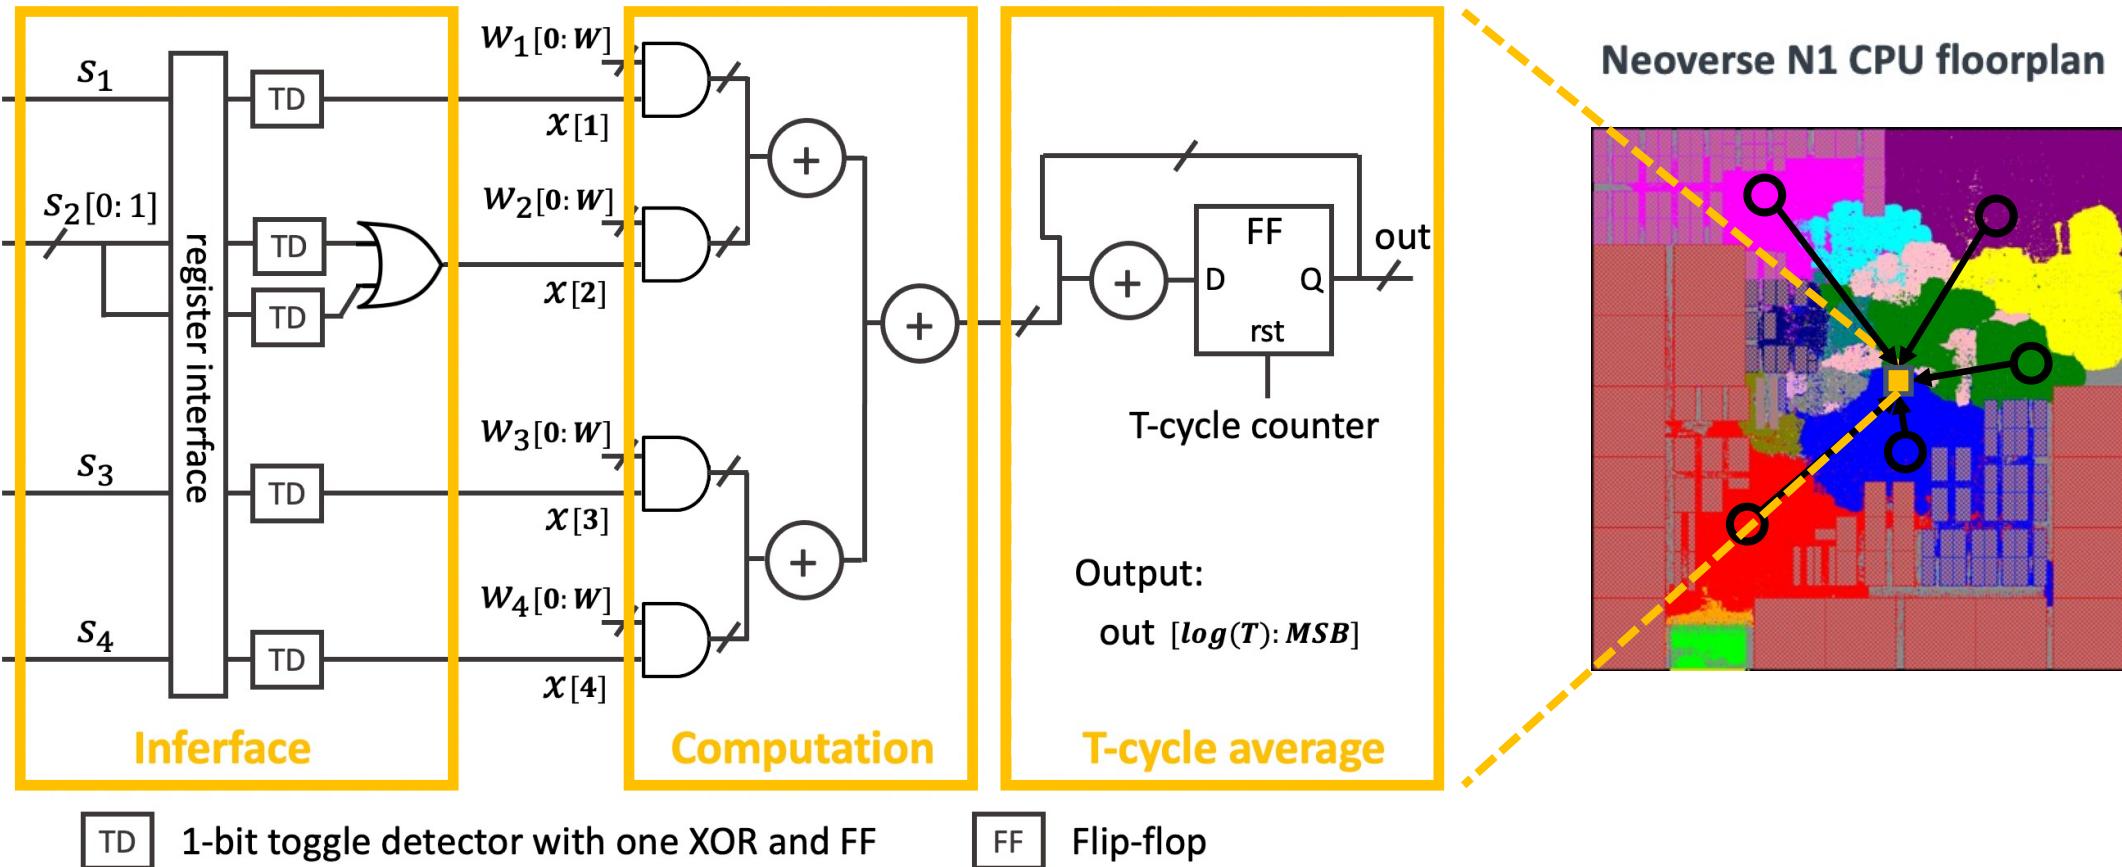
Negligible difference

Prediction from runtime OPM with $Q=159$ proxies



- MAE = 7.19%
- $R^2 = 0.953$
- $W=11$ bits after quantization


Overview of the OPM Hardware Design



- **No** multipliers or dividers, only **Q** binary inputs and **W**-bit quantized weights

Accuracy vs. Hardware Cost (Area Overhead) of the OPM

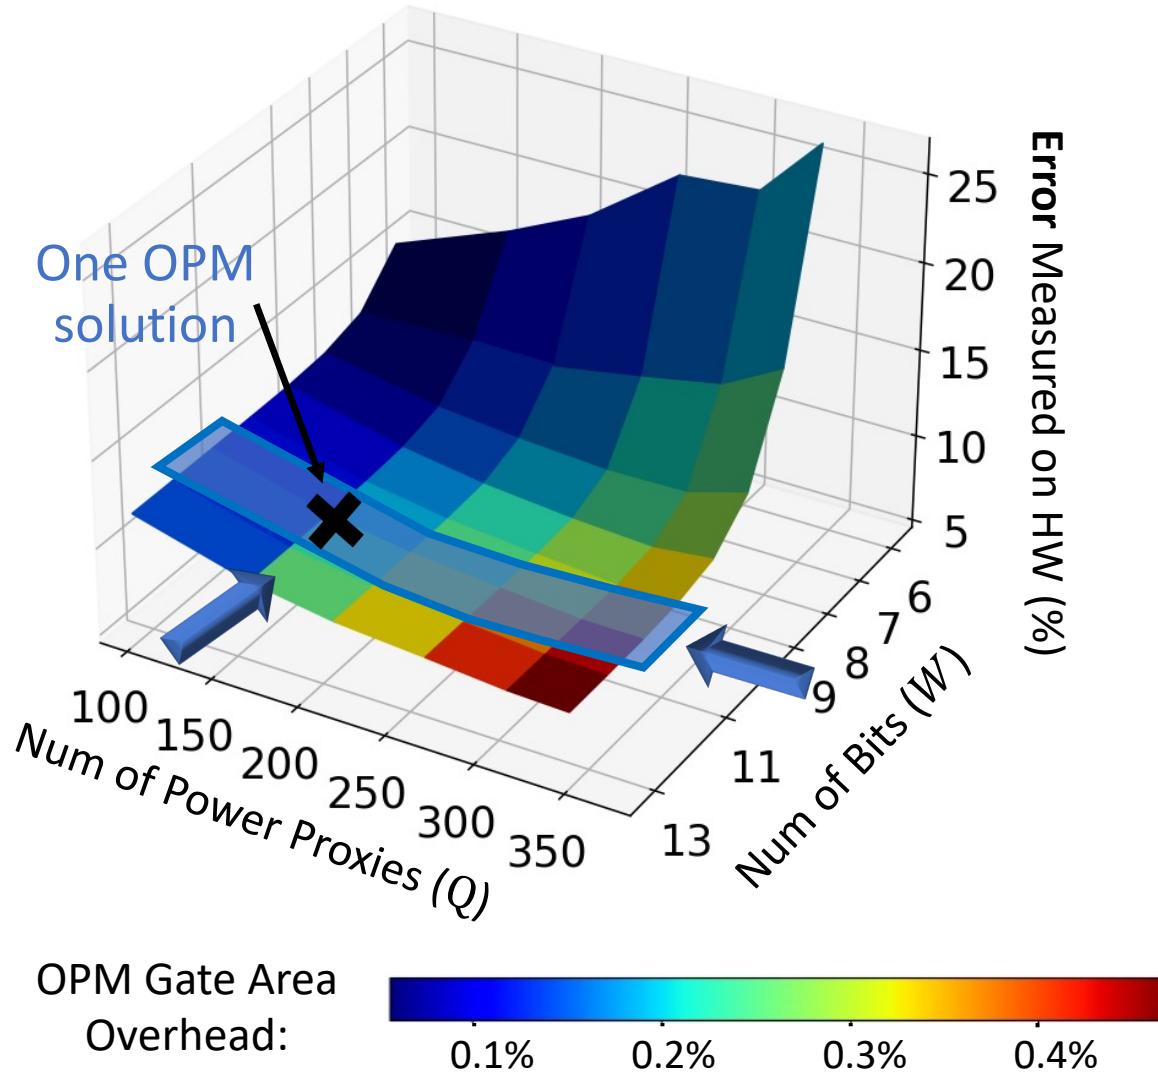
Runtime OPM implementation on Neoverse N1



- Trade-off accuracy and hardware cost
- Sweep proxy num Q and quantization bits W

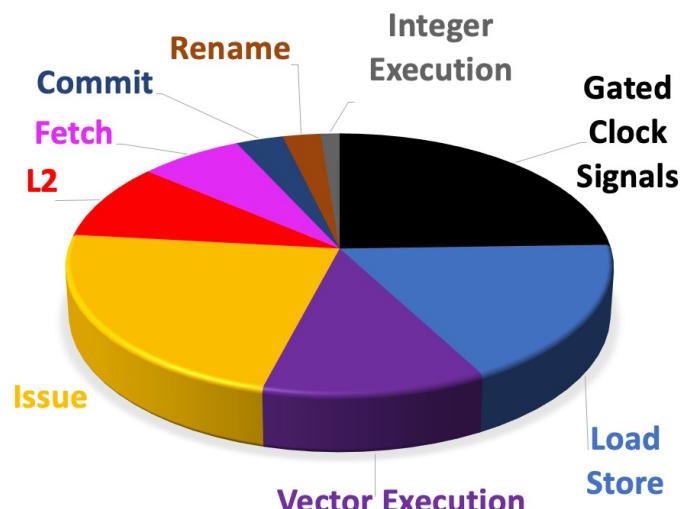
Accuracy vs. Hardware Cost (Area Overhead) of the OPM

Runtime OPM implementation on Neoverse N1

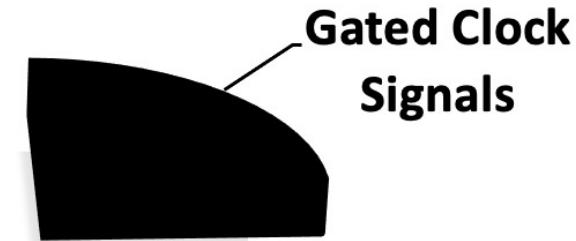


- Trade-off accuracy and hardware cost
- Sweep proxy num Q and quantization bits W
- **Strategy**
 - Keep quantization $W=10$ to 12 bits
 - Vary Q for different solutions
- **For an OPM with $Q=159$, $W=11$**
 - < 0.2% area overhead of Neoverse N1
 - < 10% in the error

Potential Application: Design-time Power Introspection



Distribution of power proxies on Neoverse N1

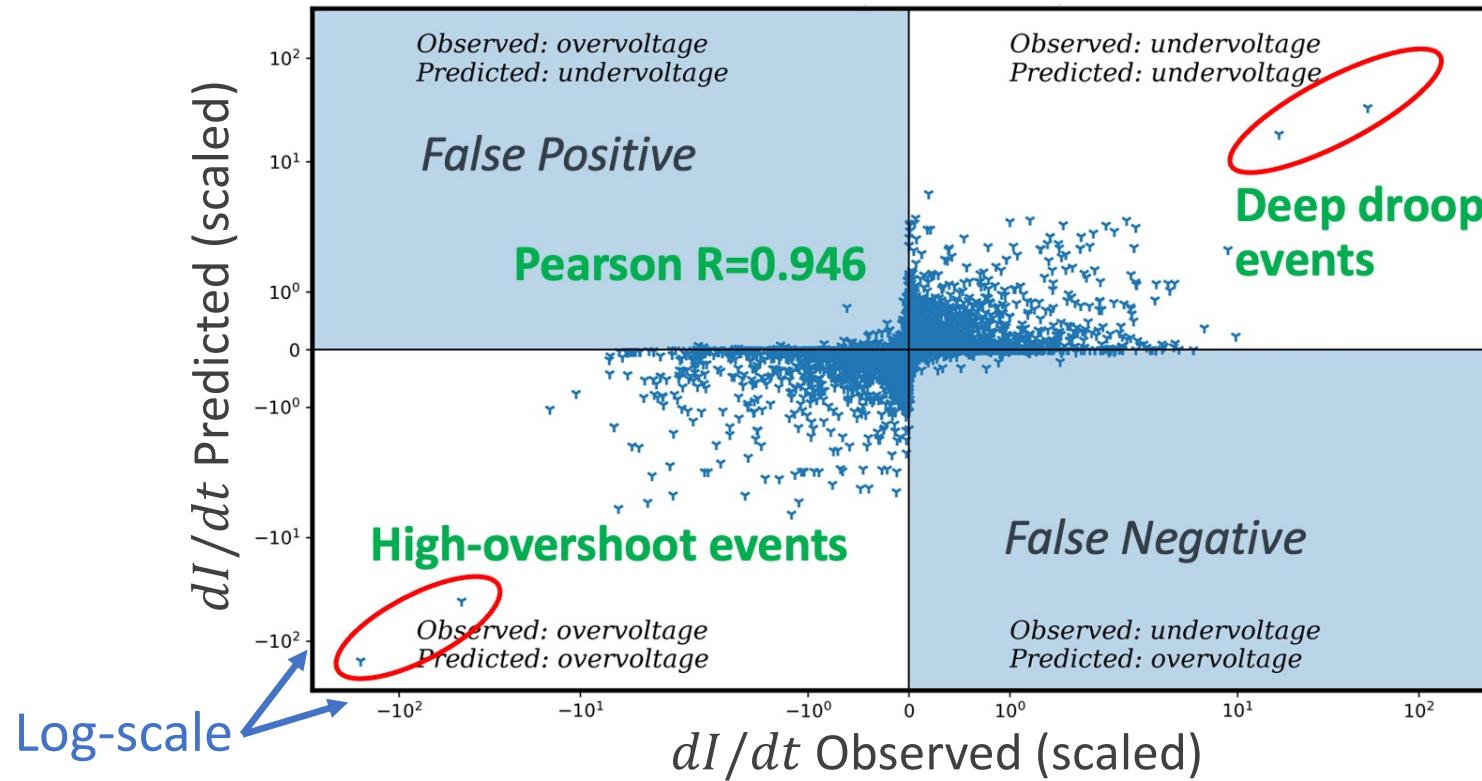


Trained **only** with more meaningful signals as initial feature candidates

Better **interpretability**


Identify power contributors for designers!

Potential Application: Runtime dI/dt Mitigation



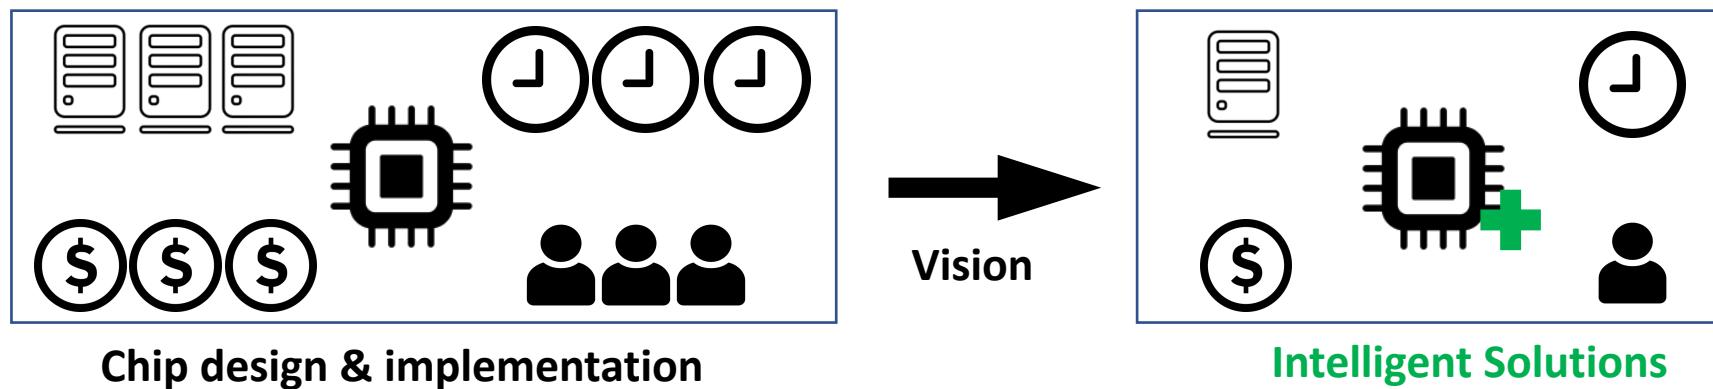
Enable CPU-driven Proactive dI/dt Mitigation!

What I Believe We Should Target




Summary and Takeaway

- Problem: Increasing Challenges in Chip Design
 - Cost, time-to-market, reliance on designers, diminishing performance return,
- **ML** in chip design
 - Less simulation time, faster feedback, less designer effort
- **AutoML** in chip design
 - Reduces months of model development to hours, no developers
- **Unified ML** in both design & runtime
 - Benefit the entire chip life cycle



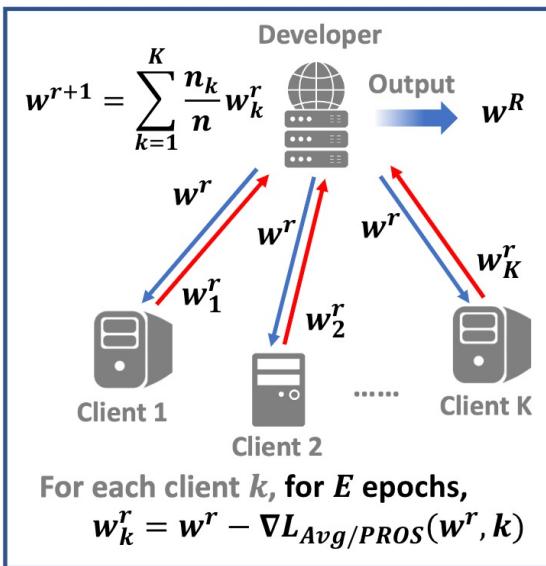
Future Research Plan



Future Works: Collaborative Framework

Collaborative ML in Chip Design

- Model quality depends on data
- Circuit data from different companies
- Design data is highly confidential



Federated Learning:
Train on local data
Communicate weights


Example – Collaborative Training

	Test on 9 Clients (C1 to C9)									Avg
	C1	C2	C3	C4	C5	C6	C7	C8	C9	
Train on 9 Clients	0.68	0.59	0.59	0.58	0.58	0.56	0.65	0.60	0.52	0.59
	0.49	0.52	0.50	0.51	0.52	0.50	0.53	0.52	0.37	0.50
	0.55	0.56	0.55	0.50	0.52	0.46	0.57	0.57	0.49	0.53
	0.52	0.49	0.51	0.53	0.51	0.53	0.52	0.52	0.46	0.51
	0.71	0.53	0.59	0.55	0.55	0.61	0.60	0.47	0.80	0.60
	0.71	0.51	0.57	0.51	0.52	0.58	0.68	0.60	0.78	0.61
	0.73	0.54	0.62	0.56	0.47	0.52	0.72	0.61	0.72	0.61
	0.76	0.60	0.65	0.60	0.55	0.55	0.71	0.64	0.57	0.63
	0.73	0.54	0.65	0.59	0.50	0.61	0.73	0.61	0.91	0.65
Train & Test Same Client	0.68	0.52	0.55	0.53	0.55	0.58	0.72	0.64	0.91	0.63
FedProx	0.63	0.83	0.71	0.72	0.66	0.67	0.63	0.57	0.42	0.65
FedProx + Finetuning	0.83	0.86	0.76	0.75	0.74	0.75	0.81	0.72	0.90	0.79

One same model in a row Nine different models in a row

- Assuming data distributed to 9 clients (C1 to C9)


Future Research Plan



Future Works: Fully-Automated & Reliable Framework

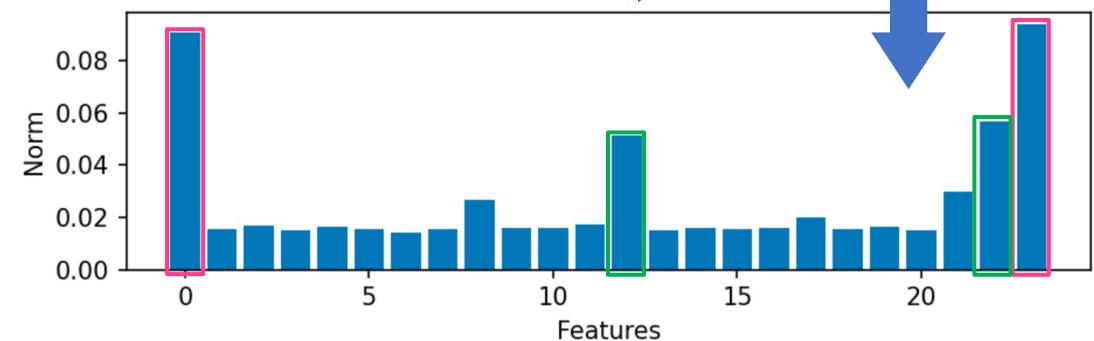
Fully-Auto ML in Chip Design

- Automated feature selection
- Automated data selection
- Automated data augmentation



Example – Feature Selection

- The first conv layer:
32 kernels with size $3 \times 3 \times 24$ 24 feature weight vectors



# Features	ROC-AUC
2 (■)	0.866
4 (■ + ■)	0.867
24 (all)	0.867

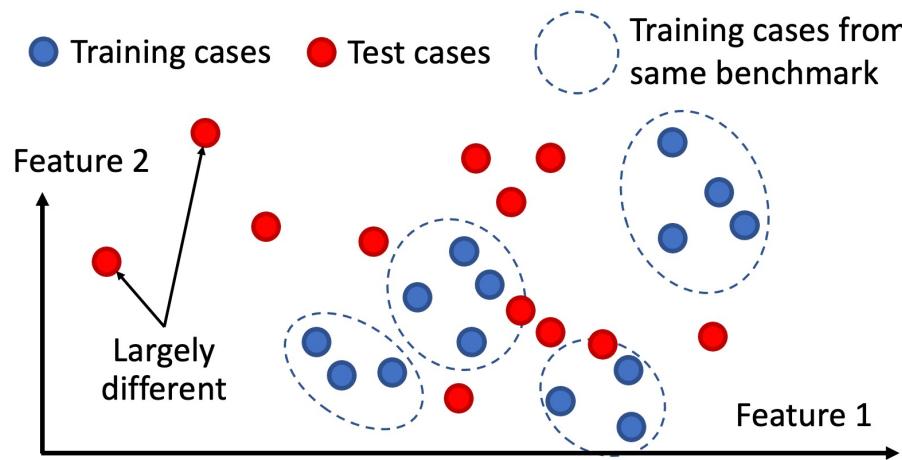
Selected Features:

0. Pin density
12. MST fly lines
22. DFF cell density
23. Clock tree cell density

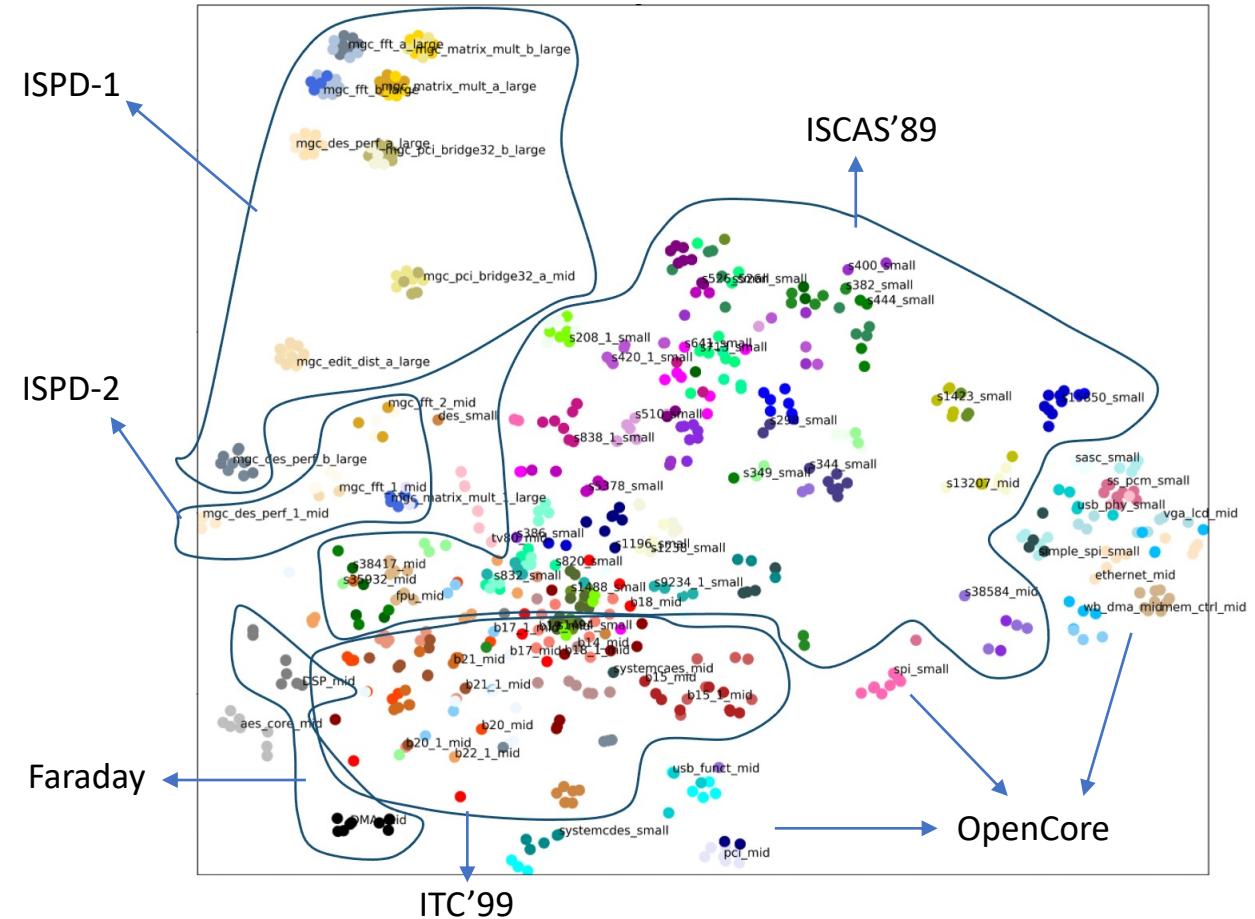
Future Works: Fully-Automated & Reliable Framework

Reliable ML in Chip Design


- Designs very sparsely distributed
- Almost impossible to perform well on every test case
- How can we trust each prediction?



Example – Design Difference



Future Research Plan



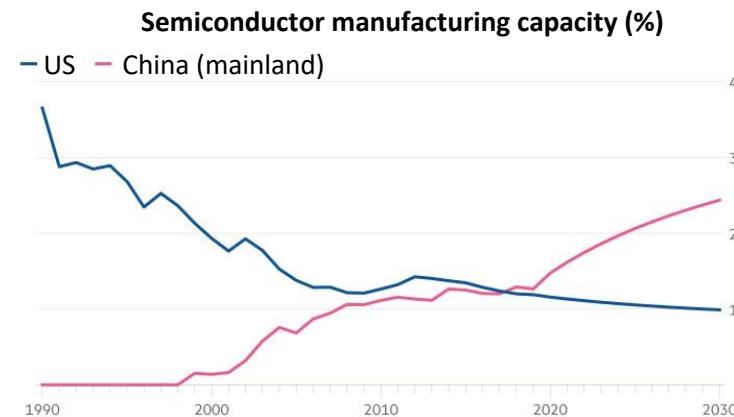
Future Funding and Collaboration Opportunities

- Agencies:
 - General Research Fund (GRF), Early Career Scheme, NSFC, ITF
- US companies:
 - Cadence, Synopsys, Nvidia, Arm, NXP
- Chinese companies:
 - Huawei, Alibaba T-head, Chinese EDA start-ups like UniVista

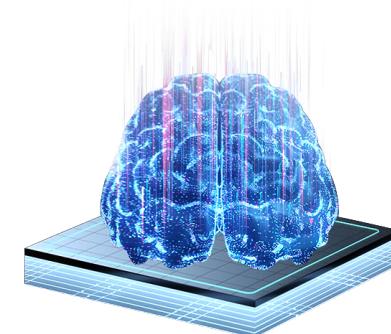


**China's New Semiconductor Policies:
Issues for Congress**

**US restricts software exports to
Chinese chip companies**



**Semiconductor switching to Asia,
including 'Greater Bay Area'**



**A great chance to overtake
leading EDA companies**

Previous Collaborations and Grant Writing Experiences

- Many thanks for my advisors and collaborators:

Prof. **Yiran Chen**
Duke University

Prof. **Hai “Helen” Li**
Duke

Prof. **Jiang Hu**
TAMU

Dr. **Brucek Khailany**
Nvidia

Dr. **Haoxing Ren**
Nvidia

Dr. **Shidhartha Das**
Arm

Dr. **Xiaoqing Xu**
Arm

Dr. **Brian Cline**
Arm

Dr. **Chand Kashyap**
Cadence

Dr. **Aiqun Cao**
Synopsys

- My previous grant writing experiences (funded):
 - **NSF**: Revitalizing EDA from a Machine Learning Perspective
 - **SRC**: A Machine Learning Approach for Cross-Level Optimizations
 - **SRC**: A Collaborative Machine Learning Approach to Fast and High-Fidelity Design Prediction
 - **Industry (Cadence)**: NAS-based Fully Automatic ML Estimator Development Flow in EDA
 - **Industry (Cadence)**: A Machine-Learning based Pre-placement Wirelength Estimator

Courses I am Qualified to Teach

- Circuit and Computer Engineering Courses
 - Digital VLSI design, digital integrated circuits (**CAD admin at Duke**)
 - Chip design methodologies
 - Digital logic & systems (**TA of undergraduate course at Duke**)
 - Computer organization and architecture
- Machine Learning Courses
 - Linear algebra for engineering (**TA of graduate course at Duke**)
 - Data mining, artificial intelligence, machine learning
 - Computer vision, deep learning

Thanks! Questions?

If you have further questions, please contact me:

zhiyao.xie@duke.edu